



Universidad de Oviedo



# Distributed and big data systems



SOFTWARE  
ARCHITECTURE

Course 2022/23

Jose E. Labra Gayo

# Distributed systems

Integration styles

Topologies: Hub % Spoke, Bus

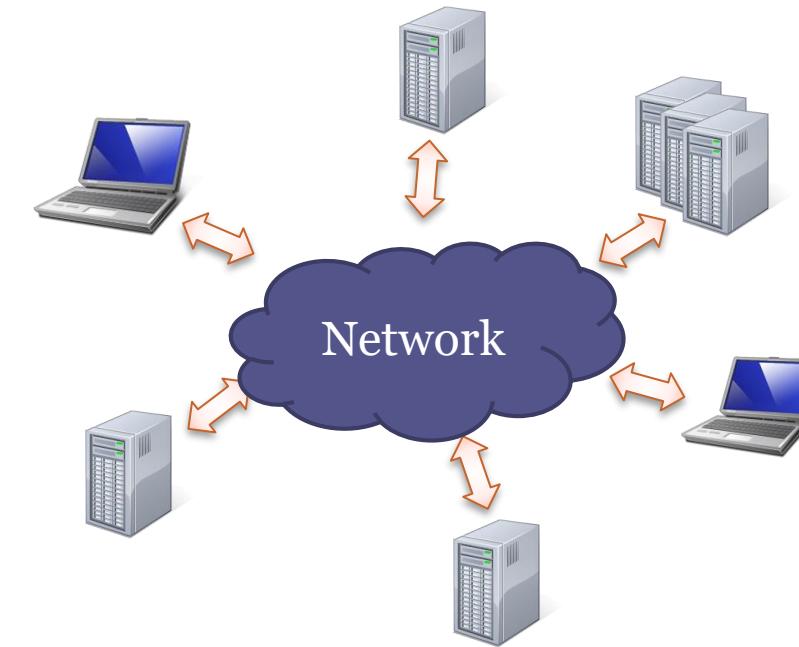
Broker pattern

Peer-to-peer

SOA

Microservices

Serverless



## Integration styles

File transfer

Shared database

Remote procedure call

Messaging

# File transfer

An application generates a data file that is consumed by another

One of the most common solutions

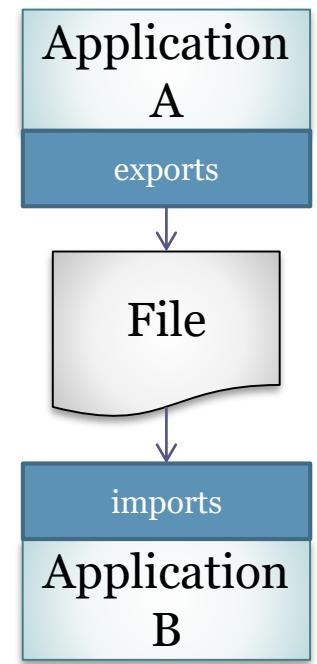
Advantages

Independence between A and B

Low coupling

Easier debugging

By checking intermediate files



# File transfer

## Challenges

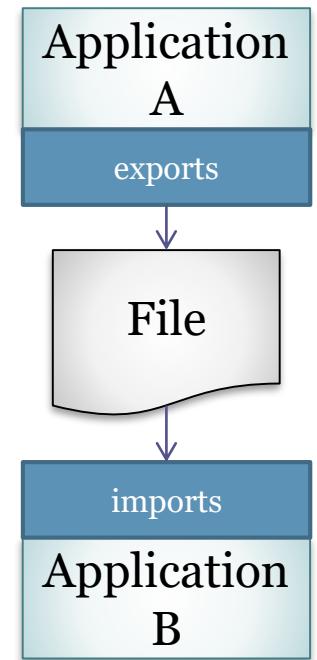
Both applications must agree a common file format

It can increase coupling

## Coordination

Once the file has been sent, the receiver could modify it  $\Rightarrow$  2 files!

It may require manual adjustments



# Shared database

Applications store their data in a shared database

Advantage

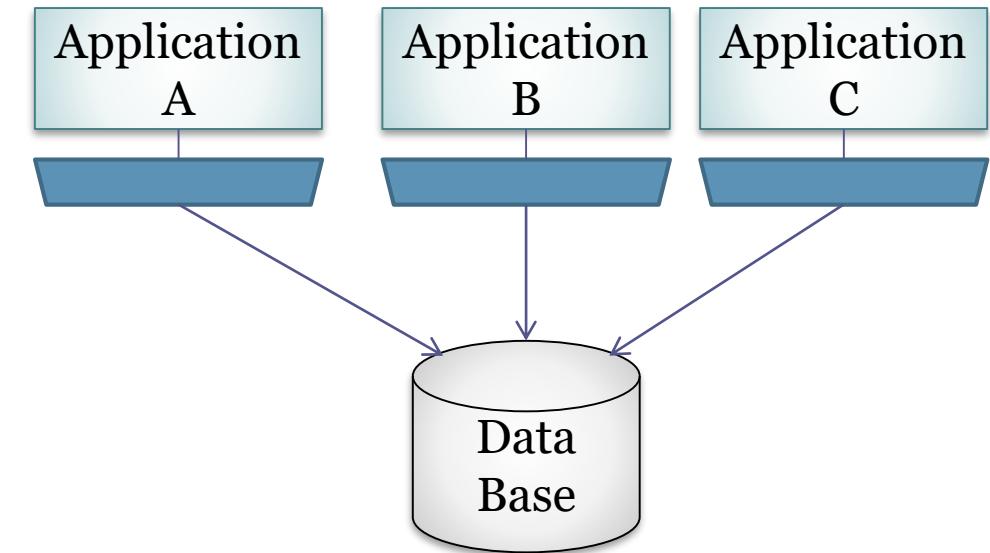
Data are always available

Everyone has access to the same information

Consistency

Familiar format

SQL for everything



# Shared database

## Challenges

### Database schema can evolve

It requires a common schema for all applications

That can cause problems/conflicts

External packages are needed (common database)

### Performance and scalability

Database as a bottleneck

### Synchronization

Distributed databases can be problematic

Scalability

NoSQL ?

# Shared database

## Variants

*Data warehousing:* Database used for data analysis and reports

ETL: process based on 3 stages

Extraction: Get data from heterogeneous sources

Transform: Process data

Load: Store data in a shared database

# Remote Procedure Call (RPC)

An application calls a function from another application that could be in another machine

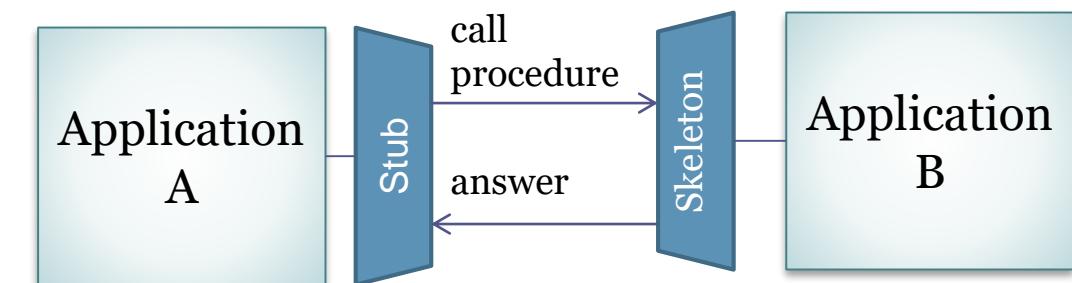
Invocation can pass parameters

Obtains an answer

Lots of applications

RPC, RMI, CORBA, .Net Remoting, ...

Web services, ...



# Remote Procedure Call (RPC)

## Advantages

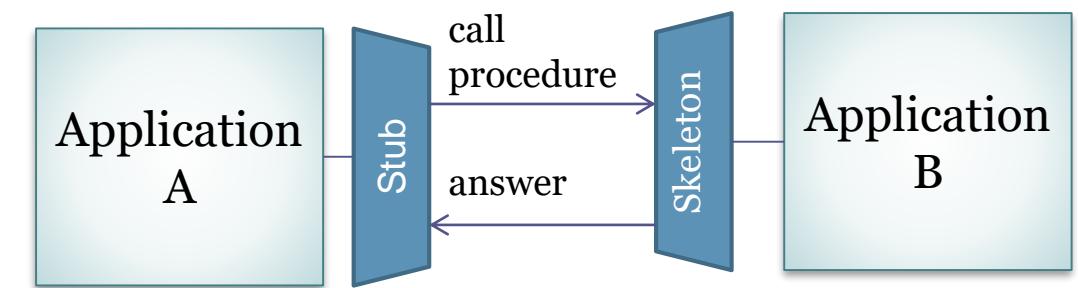
Encapsulation of implementation

Multiple interfaces for the same information

Different representations can be offered

Model familiar for developers

It is similar to invoke a method



# Remote Procedure Call (RPC)

## Challenges

### False sense of simplicity

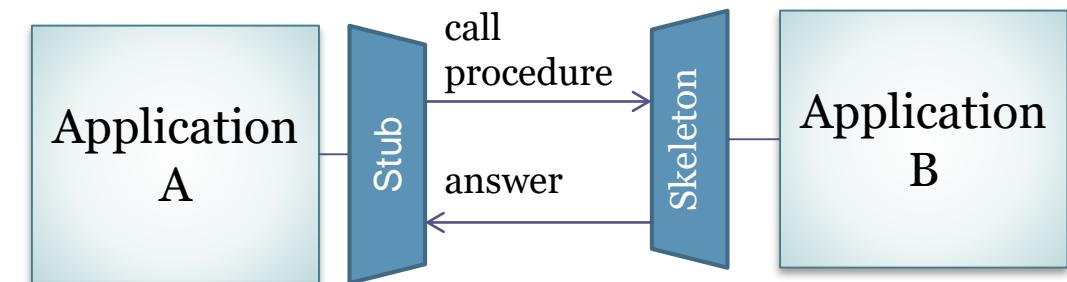
Remote procedure  $\neq$  procedure

8 fallacies of distributed computing

### Synchronous procedure calls

Increase application coupling

- The network is reliable
- Latency is zero
- Bandwidth is infinite
- The network is secure
- Topology doesn't change
- There is one administrator
- Transport cost is zero
- The network is homogeneous



# Remote procedure call

New proposals: gRPC (<https://grpc.io/>)

Google proposal

High performance RPC framework

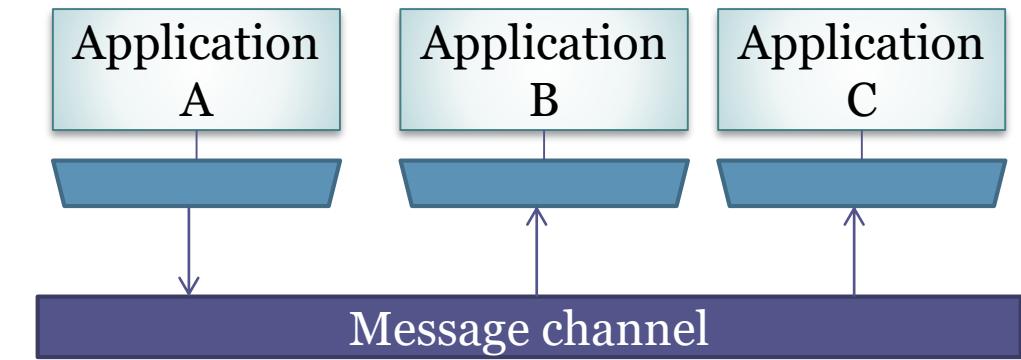
[http/2 transport protocol](#)

# Messaging

Multiple independent applications communicate sending messages through a channel

Asynchronous communication

Applications send messages and continue their execution



# Messaging

## Advantages

### Low coupling

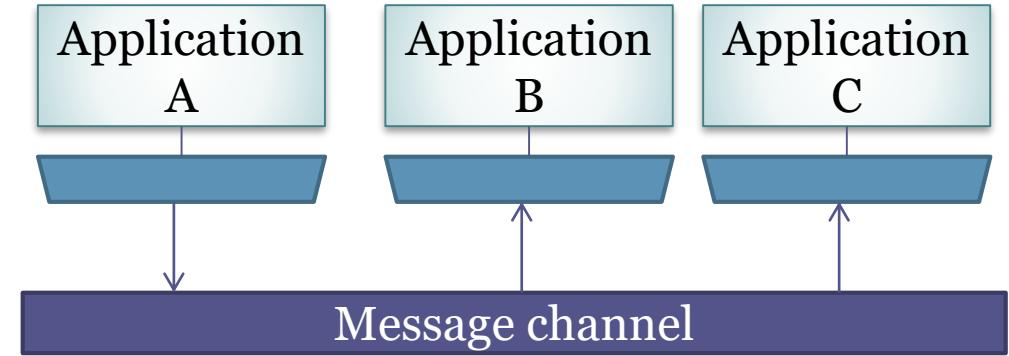
Applications are independent between each other

### Asynchronous communication

Applications continue their execution

### Implementation encapsulation

The only thing exposed is the type of messages



# Messaging

## Challenges

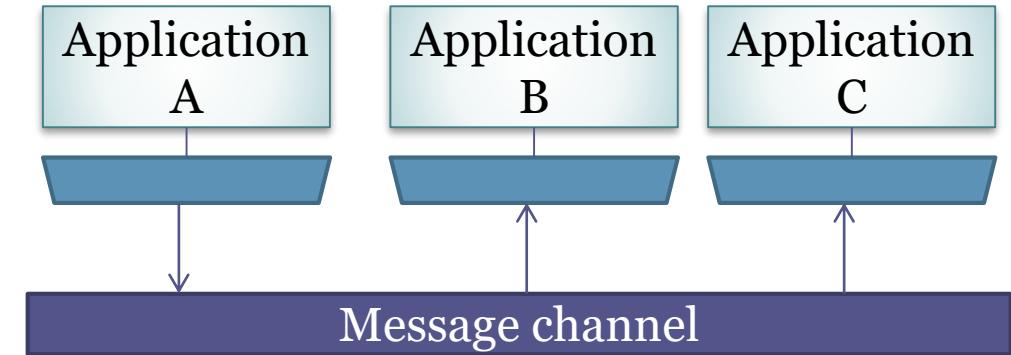
Implementation complexity

Asynchronous communication

Data transfer

Adapt message formats

Different topologies



## Integration topologies

Hub & Spoke

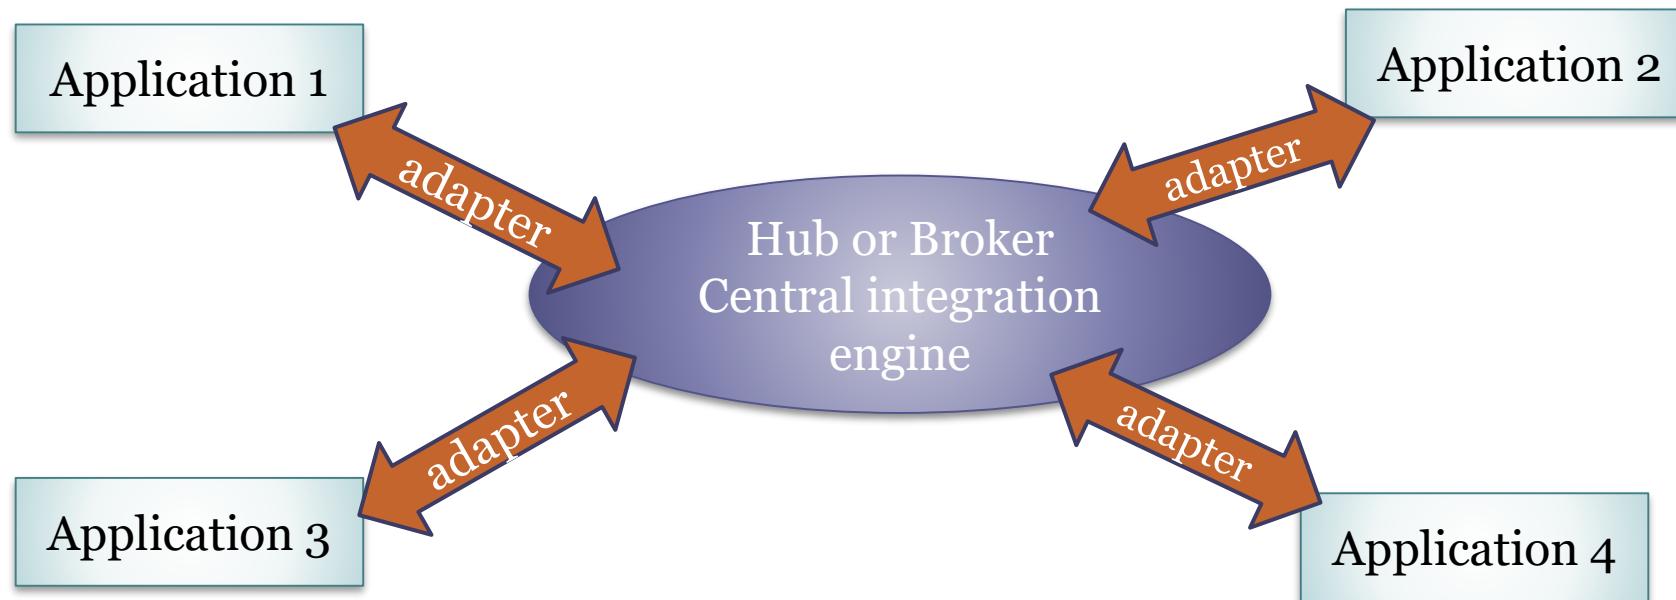
Bus

# Hub & Spoke

Related with Broker pattern

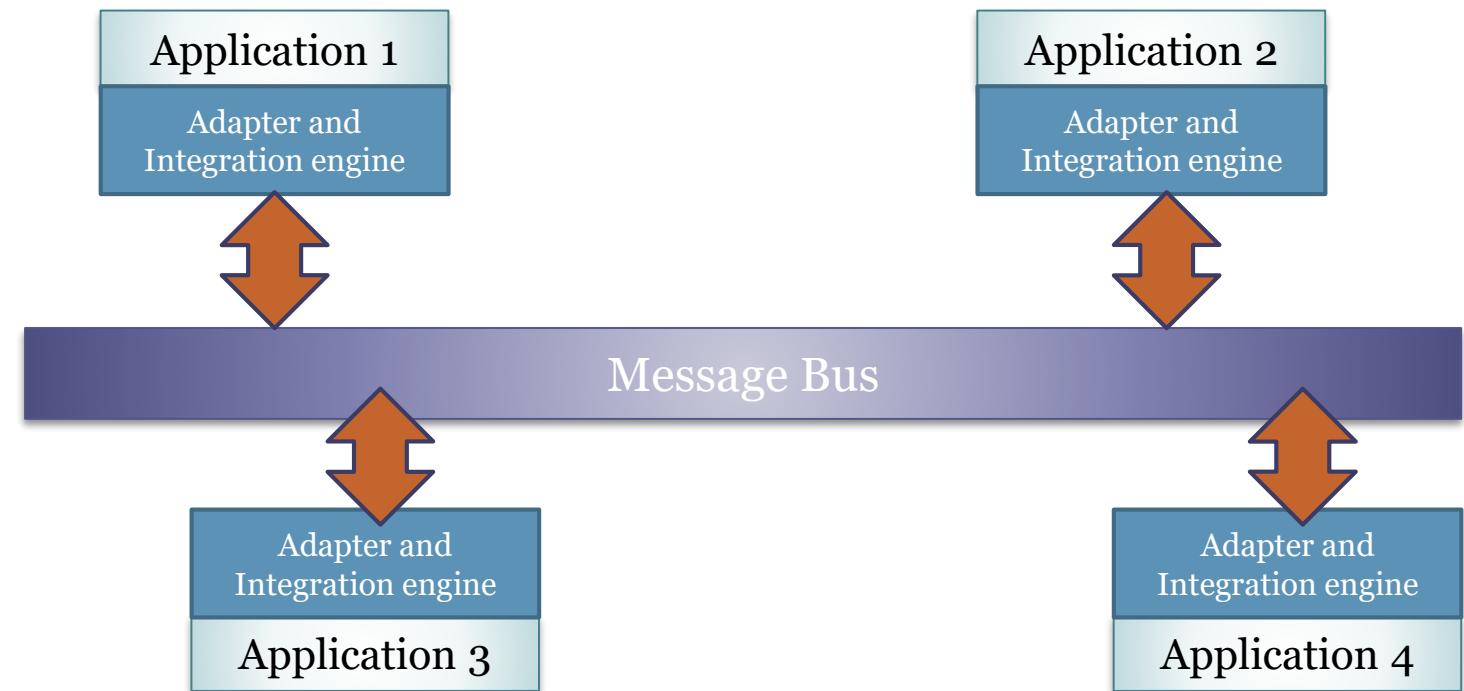
Hub = Centralized message Broker

It is in charge of integration



# Bus

Each application contains its own integration machine  
Publish/Subscribe style



# Bus

ESB - Enterprise Service Bus

Defines the messaging backbone

Some tasks

Protocol conversion

Data transformation

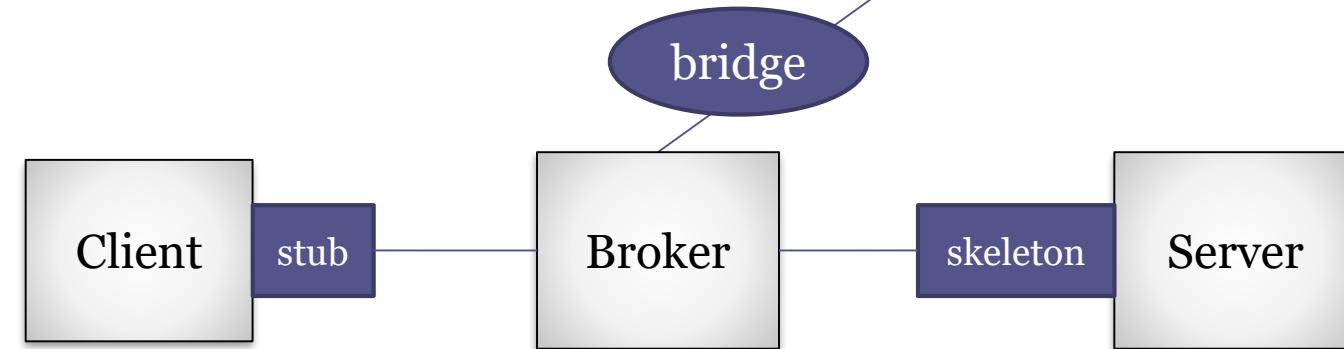
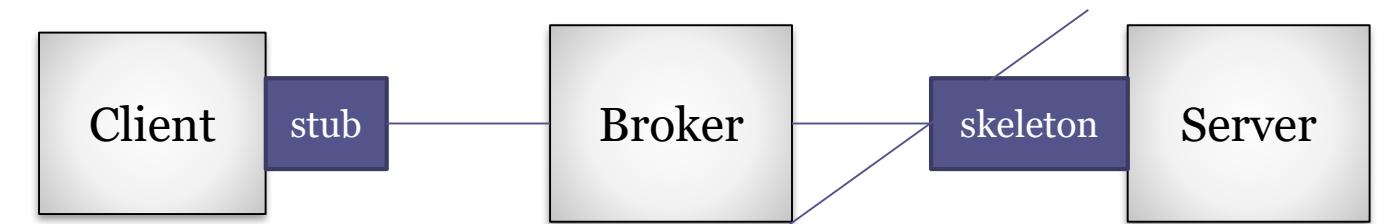
Routing

Offers an API to develop services

MOM (Message Oriented Middleware)

# Broker

Intermediate node that manages communication between a client and a server



# Broker

## Elements

### Broker

Manages communication

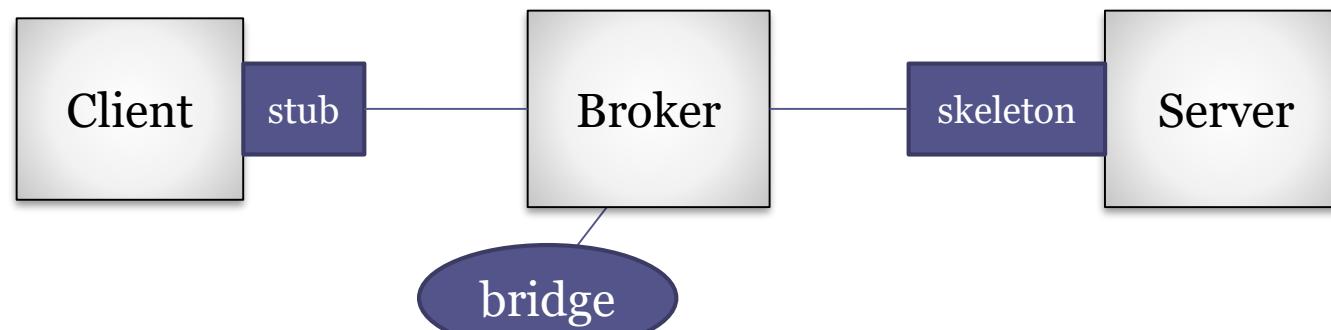
### Client: Sends requests

Client Proxy: *stub*

### Server: Returns answers

Server Proxy: *skeleton*

### Bridge: Can connect brokers



# Broker

## Advantages

**Separation of concerns**

Delegates low level communication aspects to the broker

Separate maintenance

Reusability

**Servers are independent from clients**

**Portability**

Broker = low level aspects

**Interoperability**

Using *bridges*

## Challenges

**Performance**

Adds an indirection layer

Can increase coupling between components

Broker = single point of failure

# Broker

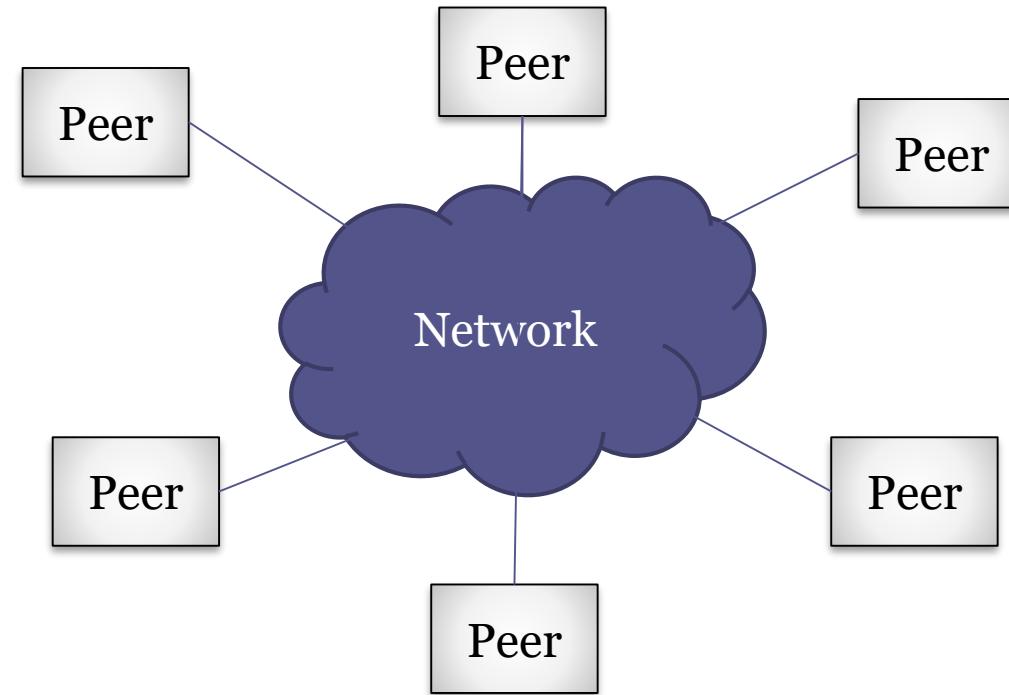
## Applications

CORBA and distributed systems

Android uses a variation of Broker pattern

# Peer-to-Peer

Equal and autonomous nodes (*peers*) that communicate between them.



# Peer-to-Peer

## Elements

Computational nodes: *peers*

They contain their own state and control thread

Network protocol

## Constraints

There is no main node

All peers are equal

# Peer-to-Peer

## Advantages

Decentralized information and control

## Fault tolerance

There is no single point of failure

A failure in one peer does not compromise the whole system

## Challenges

Keeping the state of the system

Complexity of the protocol

## Bandwidth Limitations

Network and protocol latency

## Security

Detect malicious peers

# Peer-to-Peer

Popular applications

Napster, BitTorrent, Gnutella, ...

This architecture style is not only to share files

e-Commerce (B2B)

Collaborative systems

Sensor networks

Blockchain

...

Variants

Super-peers

# Service Oriented Architectures

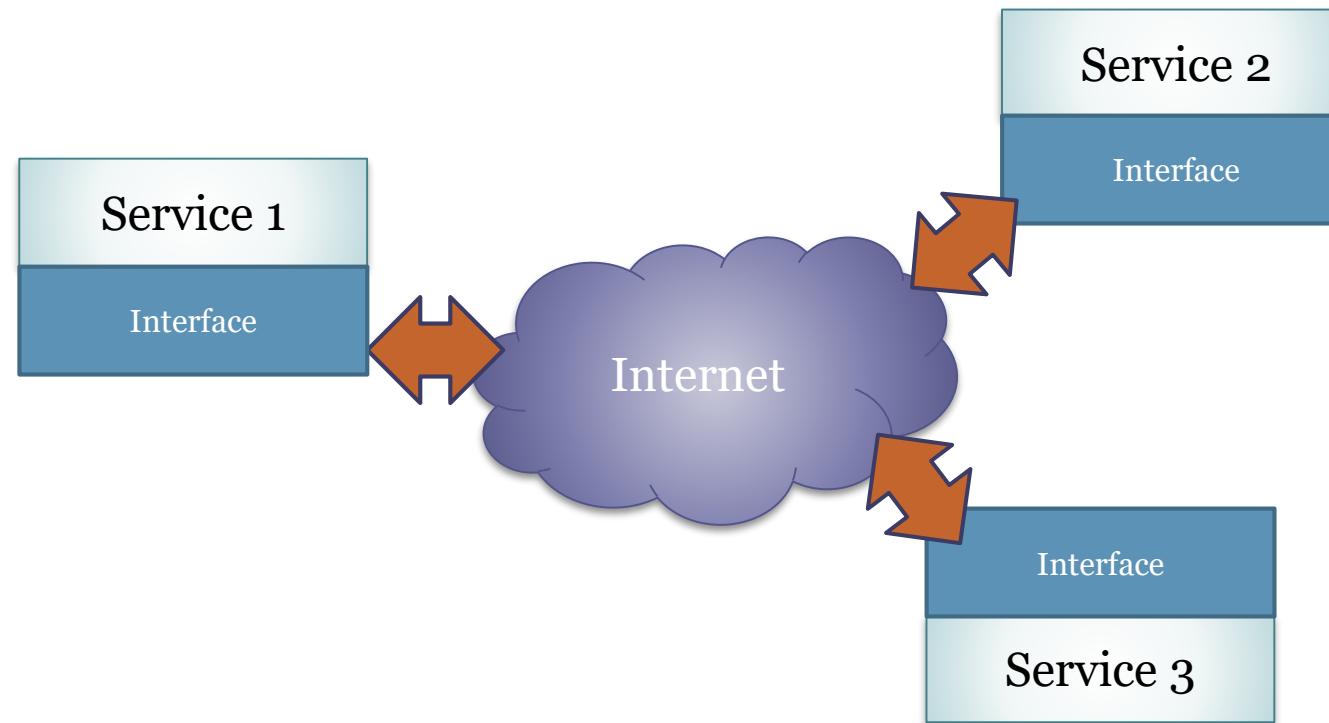
SOA

WS-\*

REST

# SOA

SOA = Service Oriented Architecture  
Services are defined by an interface



# SOA

## Elements

Provider: Provides service

Consumer: Does requests to the service

Messages: Exchanged information

Contract: Description of the functionality provided by the service

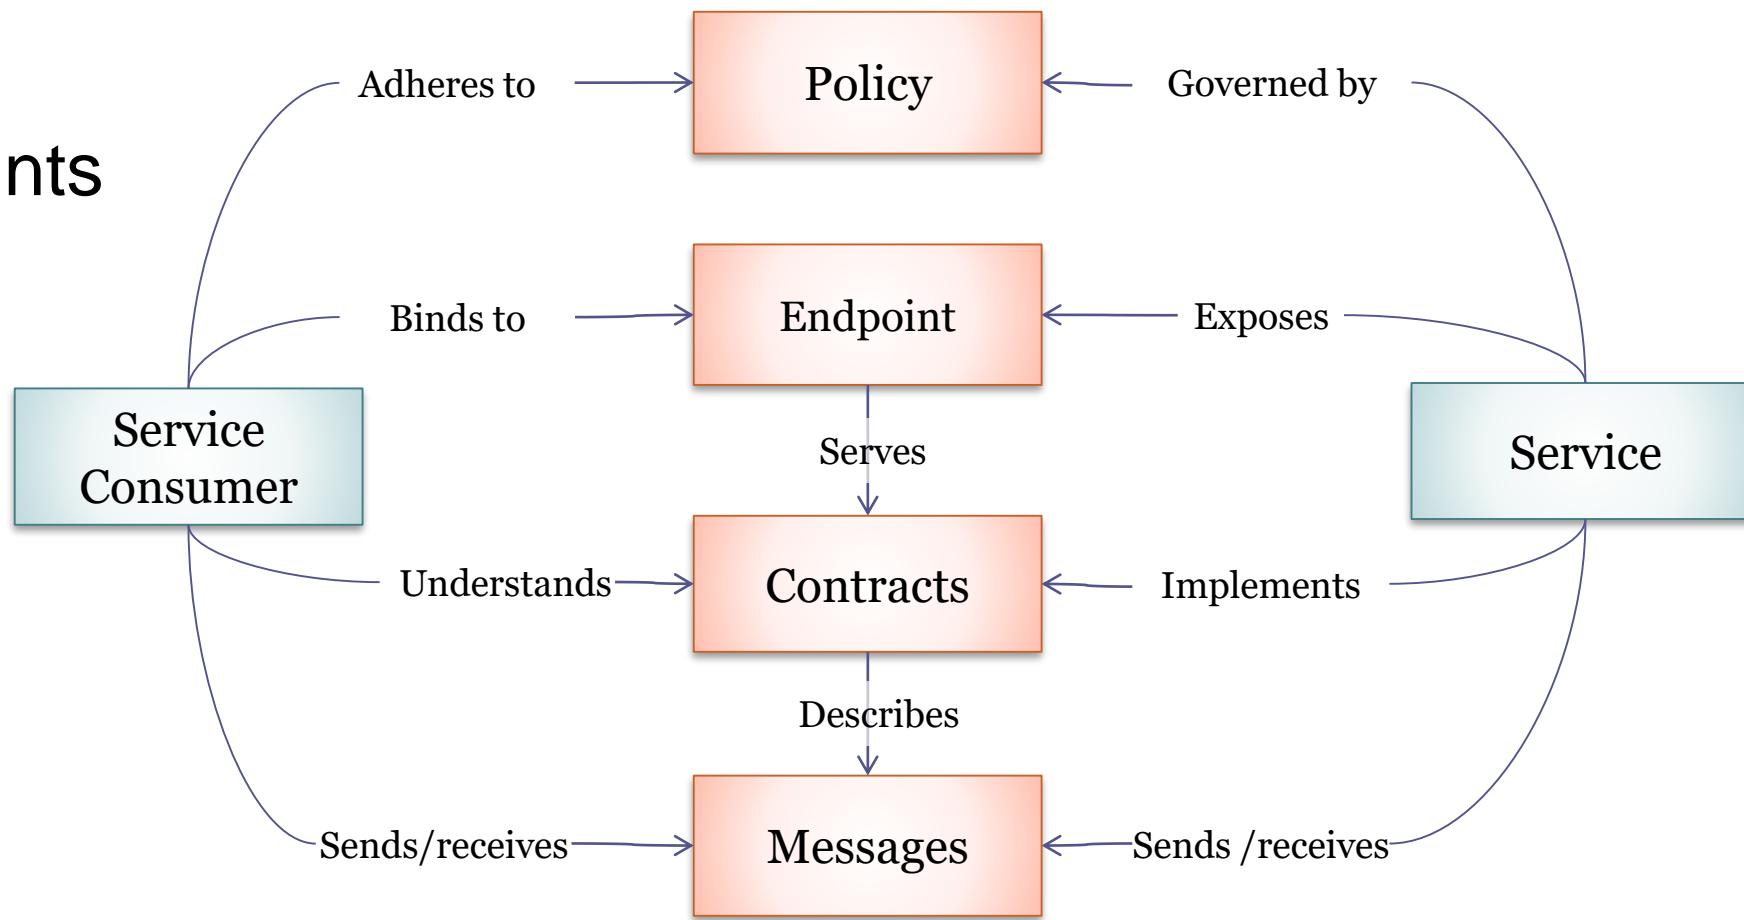
Endpoint: Service location

Policy: Service level agreements

Security, performance, etc.

# SOA

## Constraints



# SOA

## Advantages

Independent of language and platform

Interoperability

Use of standards

Low coupling

Decentralized

Reusability

Scalability

one-to-many vs one-to-one

Partial solution for legacy systems

Adding a web services layer

## Challenges

Performance

E.g. real time systems

Overkill in very homogeneous environments

Security

Risk of public exhibition of API to external parties

DoS attacks

Service composition and coordination

# SOA

Variants:

WS-\*

REST

## WS-\*

WS-\* model = Set of specifications

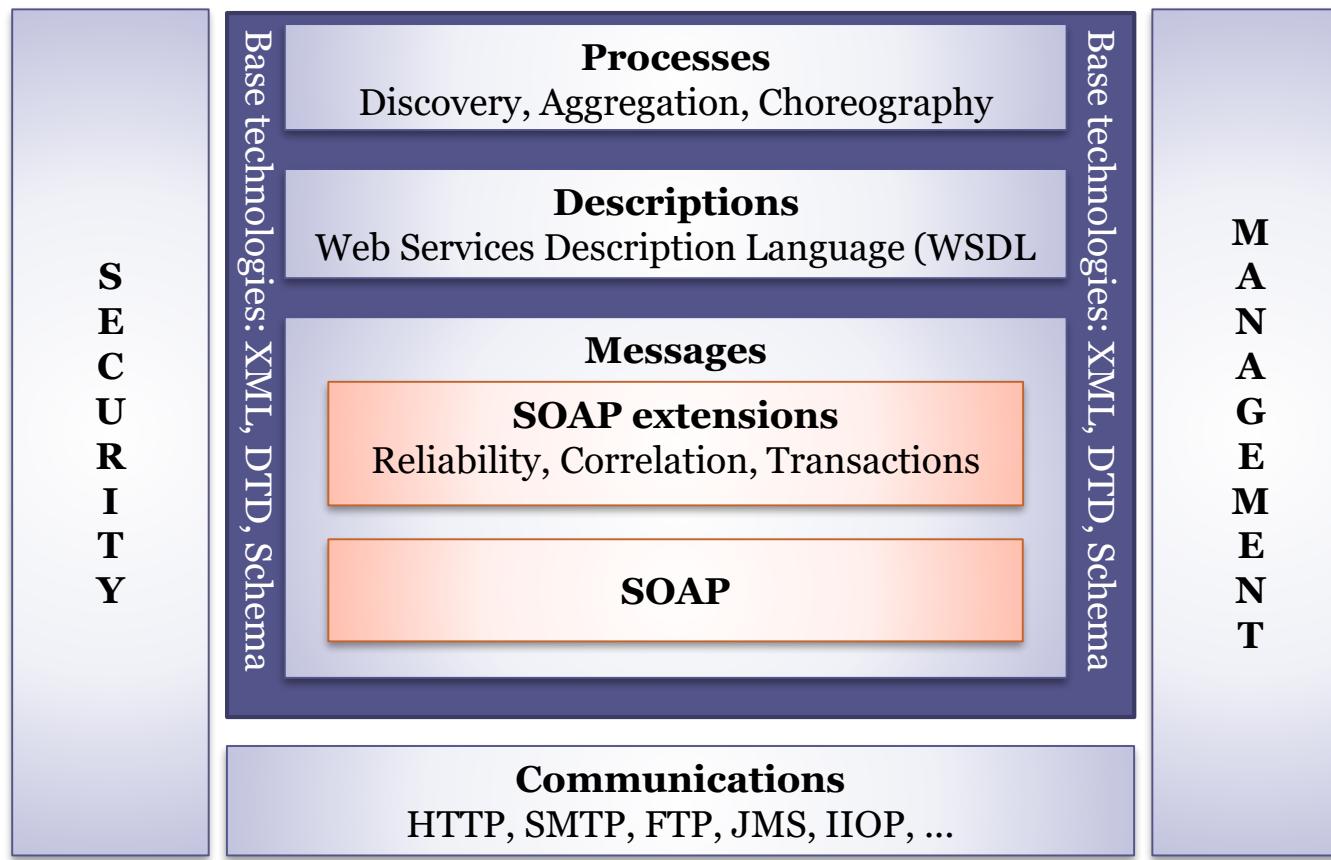
SOAP, WSDL, UDDI, etc....

Proposed by W3C, OASIS, WS-I, etc.

Goal: Reference SOA implementation

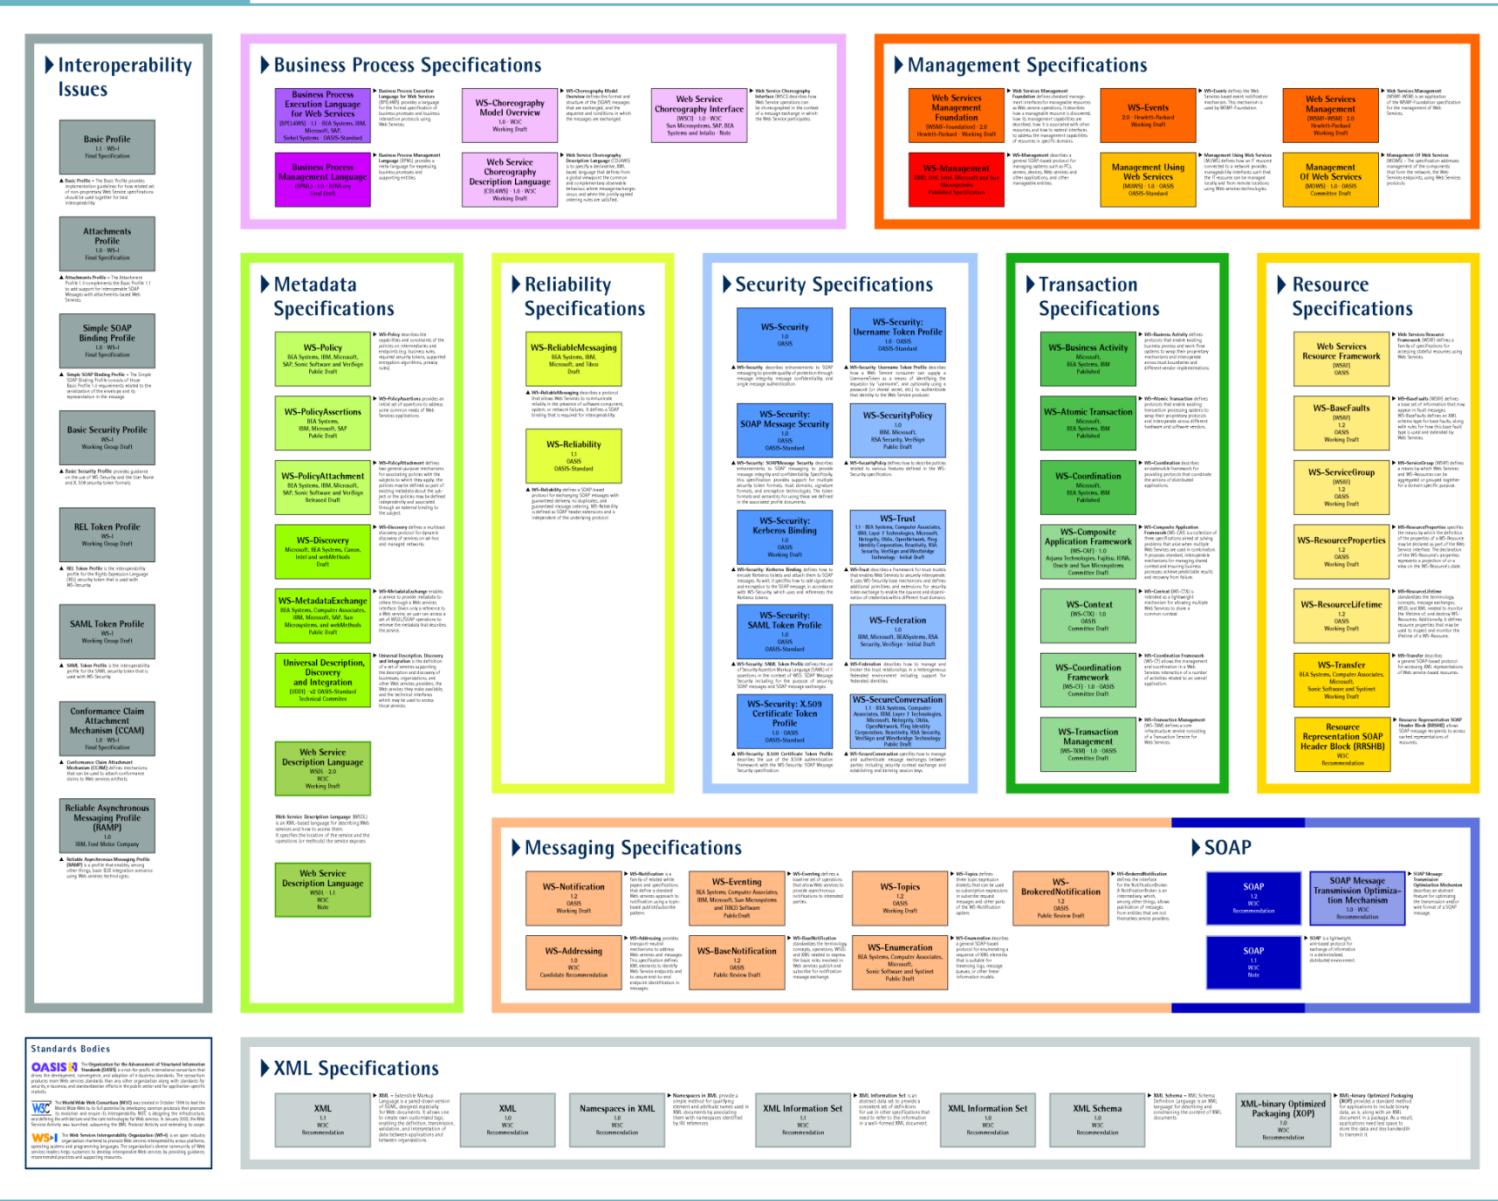
WS-\*

# Web Services Architecture



Version 2.0 - September 2005

# Web Services Standards



Deutsche Post

Deutsche Post AG  
SOP Group  
Staatschensweg 10  
53170 Bonn  
Phone +49 (220) 182 19019  
Fax +49 (220) 182 19099  
SOP-Group@DeutschePost.de  
www.SOP-Group.com

## Dependencies

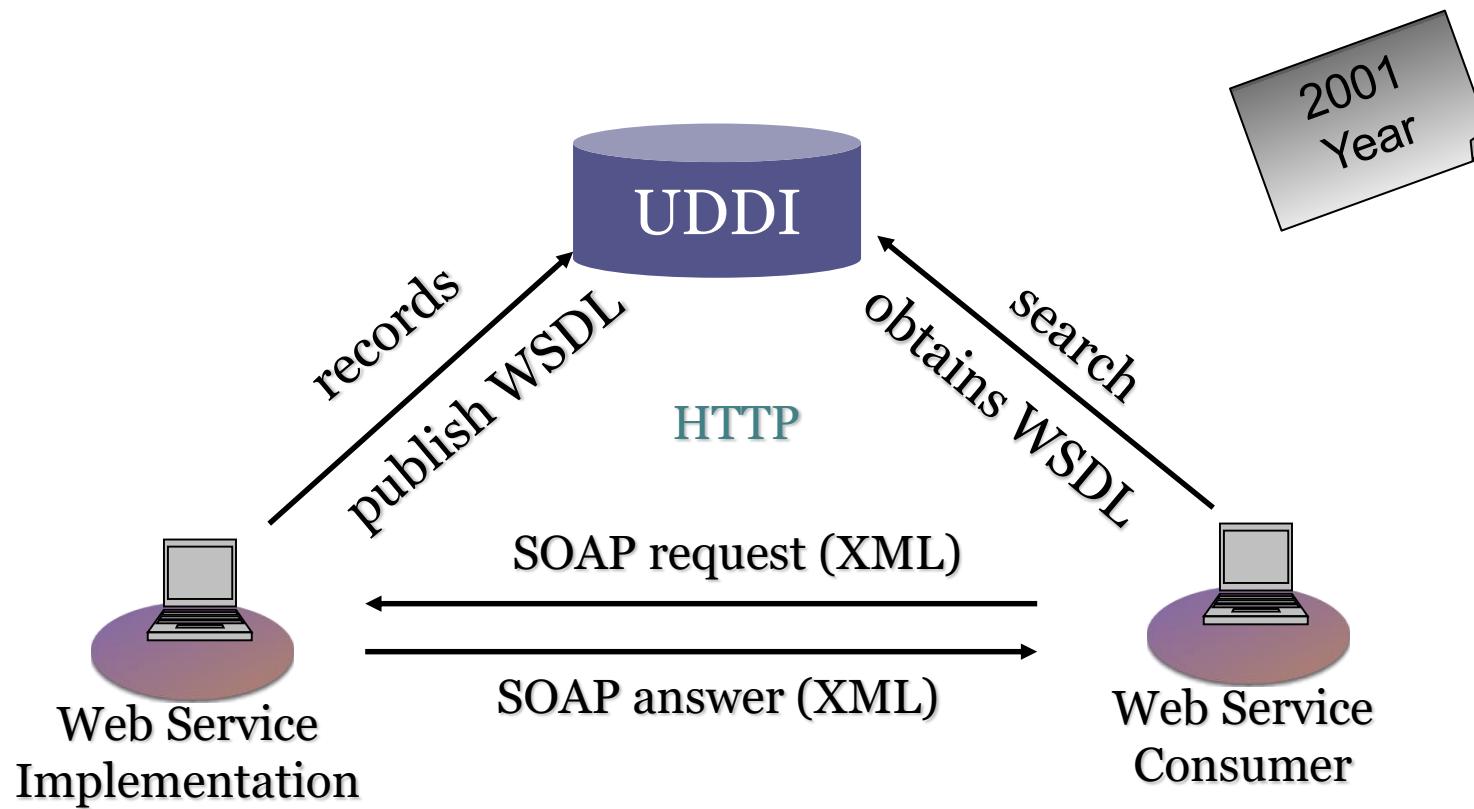


innoQ

innoQ Deutschland GmbH  
Universitätsstrasse 11  
D-40889 Ratingen  
Telefon +49 (0) 2102 - 77 16-100  
Telefax +49 (0) 2102 - 77 16-011  
Telefax +49 (0) 41 - 743 0119  
info@innog.com • www.innog.com

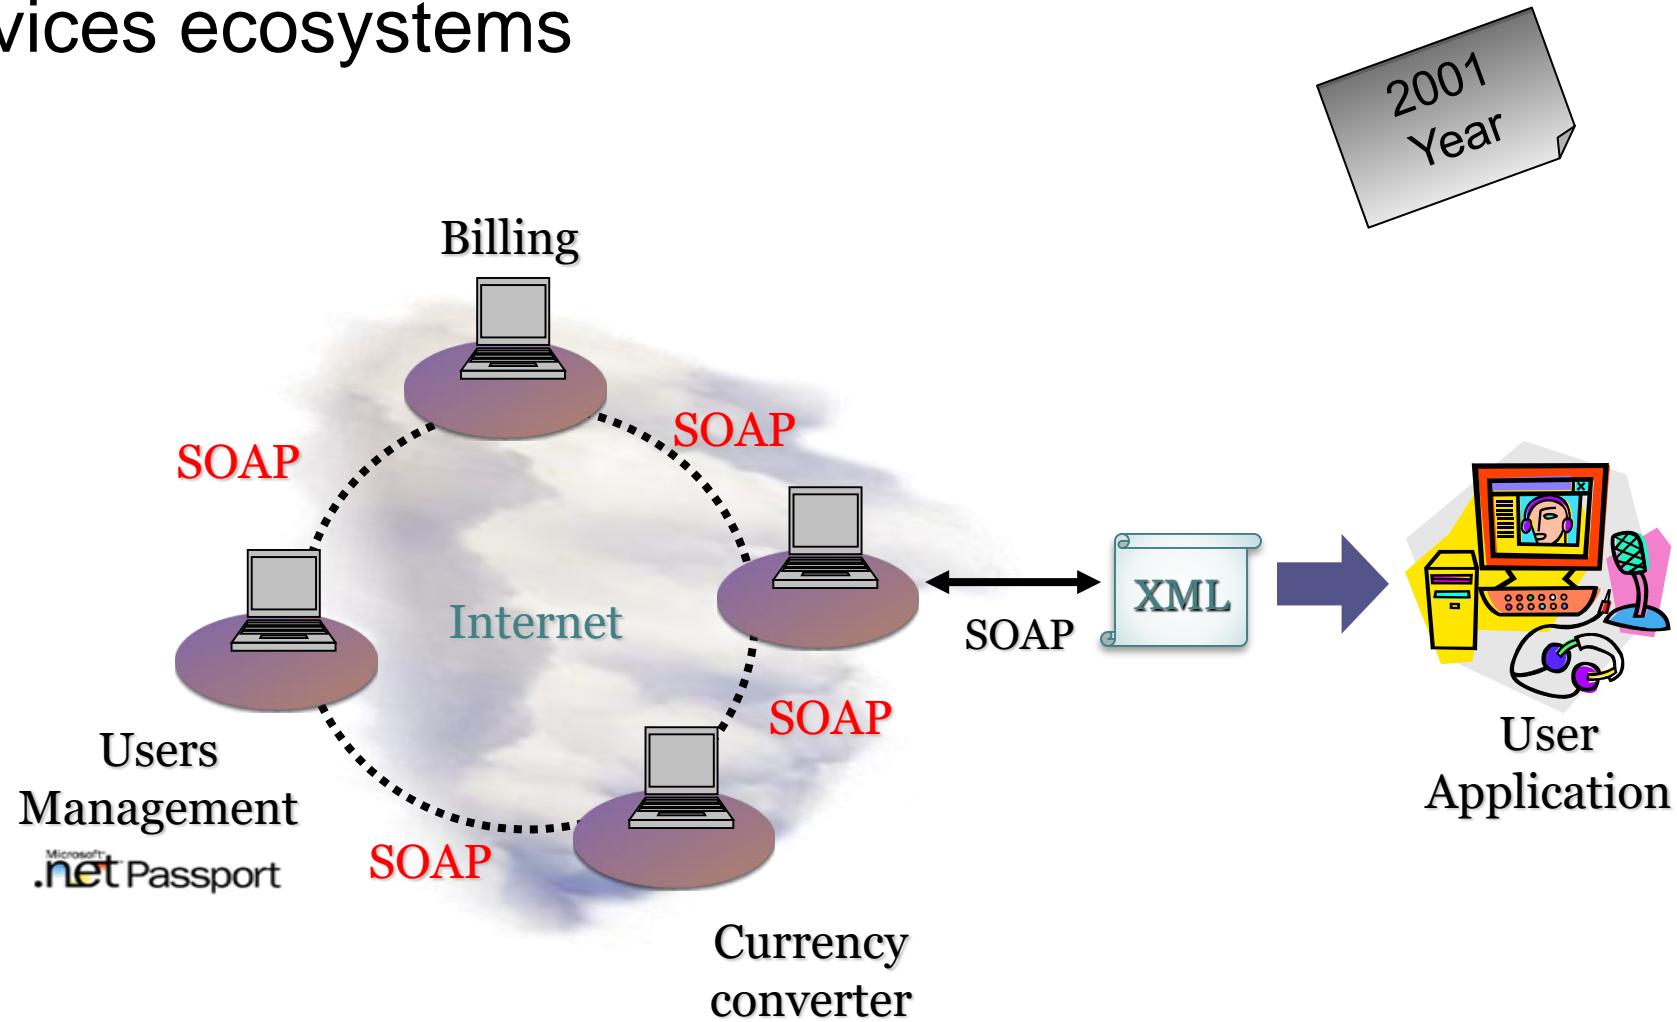
innoQ Schweiz GmbH  
Universitätsstrasse 11  
CH-6330 Cham  
Telefon +41 (0) 41 - 743 0111  
Telefax +41 (0) 41 - 743 0119

## WS-\*



## WS-\*

## Web Services ecosystems



# WS-\*

## SOAP

Defines messages format and bindings with several protocols

Initially *Simple Object Access Protocol*

### Evolution

Developed from XML-RPC

SOAP 1.0 (1999), 1.1 (2000), 1.2 (2007)

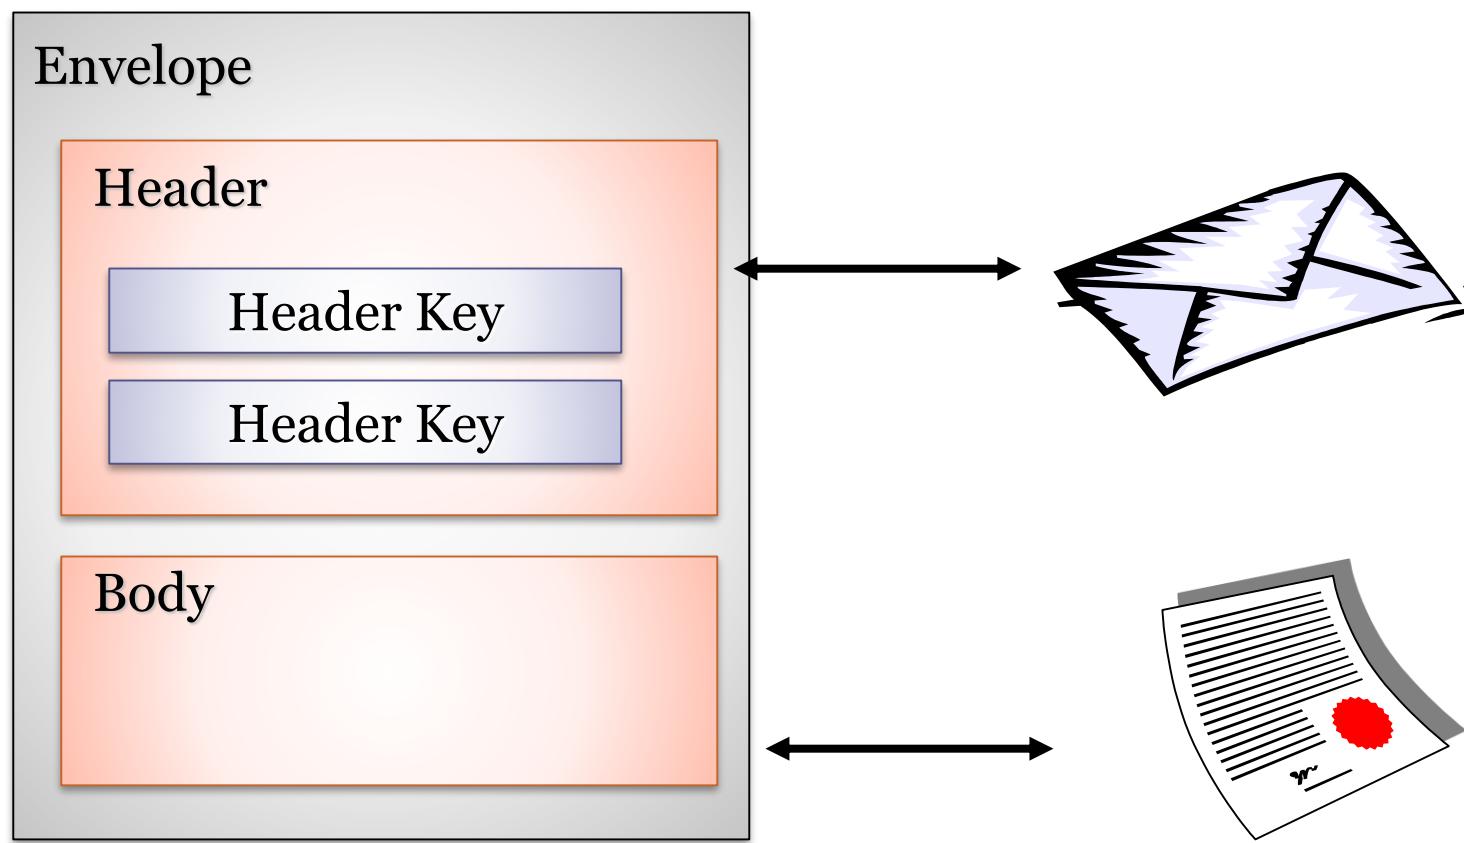
Initial development by Microsoft

Posterior adoption by IBM, Sun, etc.

Good Industrial adoption

## WS-\*

## Message format in SOAP



## WS-\*

## Example of SOAP over HTTP

2001  
Year

POST ?

```
POST /Suma/Service1.asmx HTTP/1.1
Host: localhost
Content-Type: text/xml; charset=utf-8
Content-Length: longitud del mensaje
SOAPAction: "http://tempuri.org/suma"
<?xml version="1.0" encoding="utf-8"?>
<soap:Envelope
    xmlns:soap="http://schemas.xmlsoap.org/soap/envelope/">
<soap:Body>
    <sum xmlns="http://tempuri.org/">
        <a>3</a>
        <b>2</b>
    </sum>
</soap:Body>
</soap:Envelope>
```

# WS-\*

## Advantages

Specifications developed by community

W3c, OASIS, etc.

Industrial adoption

Implementations

Integral view of web services

Numerous extensions

Security, orchestration, choreography, etc.

## Challenges

Not all specifications were mature

Over-specification

Lack of implementations

RPC style abuse

Uniform interface

Sometimes, bad use of HTTP architecture

Overload of GET/POST methods

# WS-\*

## Applications

Lots of applications have been using SOAP

Example: eBay (50mill. SOAP transactions/day)

But...some popular web services ceased to offer SOAP support

Examples: Amazon, Google, etc.

# REST

REST = REpresentational State Transfer

Architectural style

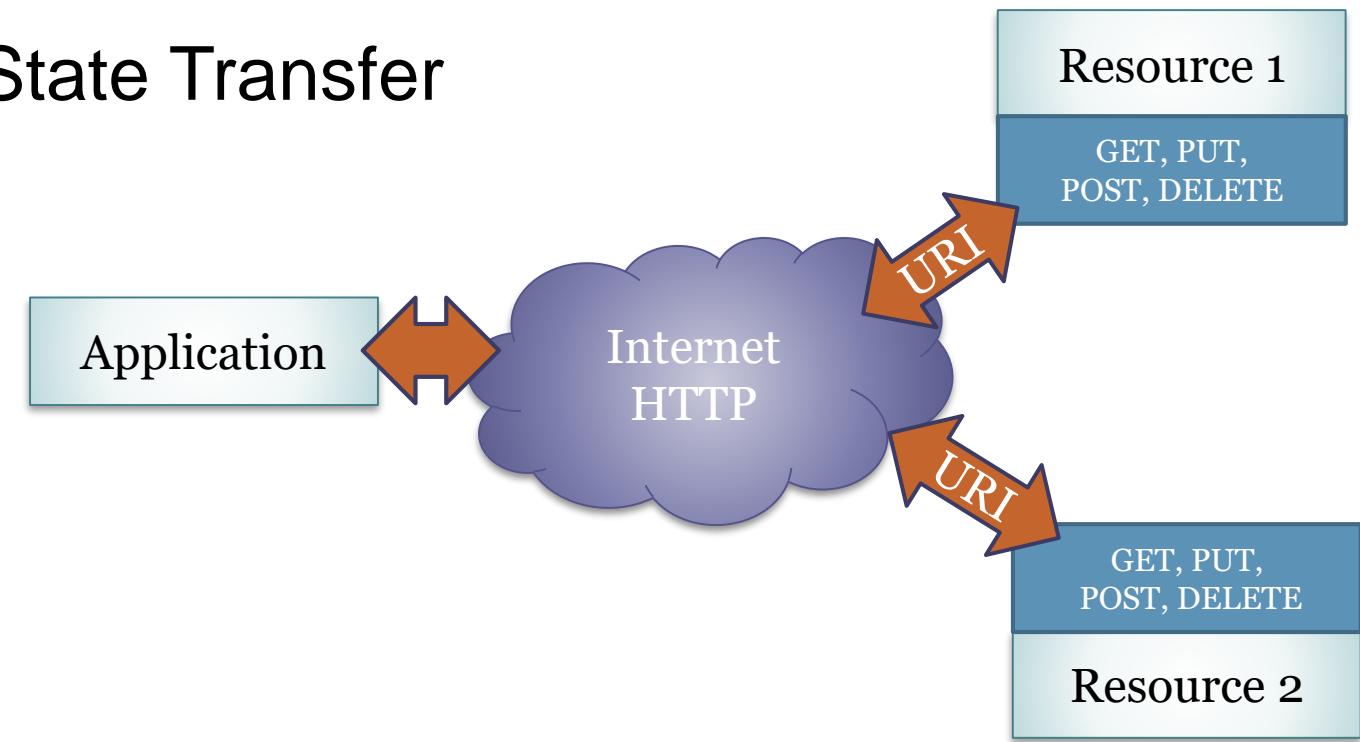
Source: Roy T Fielding PhD dissertation (2000)

Inspired by Web architecture (HTTP/1.1)



# REST

## REST - Representational State Transfer Diagram



# REST

Set of constraints

Resources with uniform interface

Identified by URIs

Fixed set of actions: GET, PUT, POST, DELETE

Resource representations are returned

Stateless

REST = Architectural style

Some levels of adoption:

RESTful

REST-RPC hybrid

# REST as a composed style

Layers

Client-Server

Stateless

Cached

Replicated server

Uniform interface

Resource identifiers (URIs)

Auto-descriptive messages (MIME types)

Links to other resources (HATEOAS)

Code on demand (optional)

# REST uniform interface

Fixed set of operations  
GET, PUT, POST, DELETE

Method	In databases	Function	Safe?	Idempotent?
PUT	≈Create/Update	Create/update	No	Yes
POST	≈Update	Create/ Update children	No	No
GET	Retrieve	Query resource info	Yes	Yes
DELETE	Delete	Delete resource	No	Yes

Safe = Does not modify server data

Idempotent = The effect of executing N-times is the same as executing it once

# REST

Stateless client/server protocol

State handled by client

**HATEOAS** (*Hypermedia As The Engine of Application State*)

Representations return URIs to available options

Chaining of resource requests

**Example:** Student management

1.- Get list of students

GET `http://example.org/student`

Returns list of students with each student URI

2.- Get information about an specific student

GET `http://example.org/student/id2324`

3.- Update information of an specific student

PUT `http://example.org/student/id2324`

# REST

## Advantages

### Client/Server

Separation of concerns

Low coupling

### Uniform interface

Facilitates comprehension

Independent development

### Scalability

Improves answer times

Less network load (cached)

Less bandwidth

## Challenges

### REST partially adopted

Just using JSON or XML

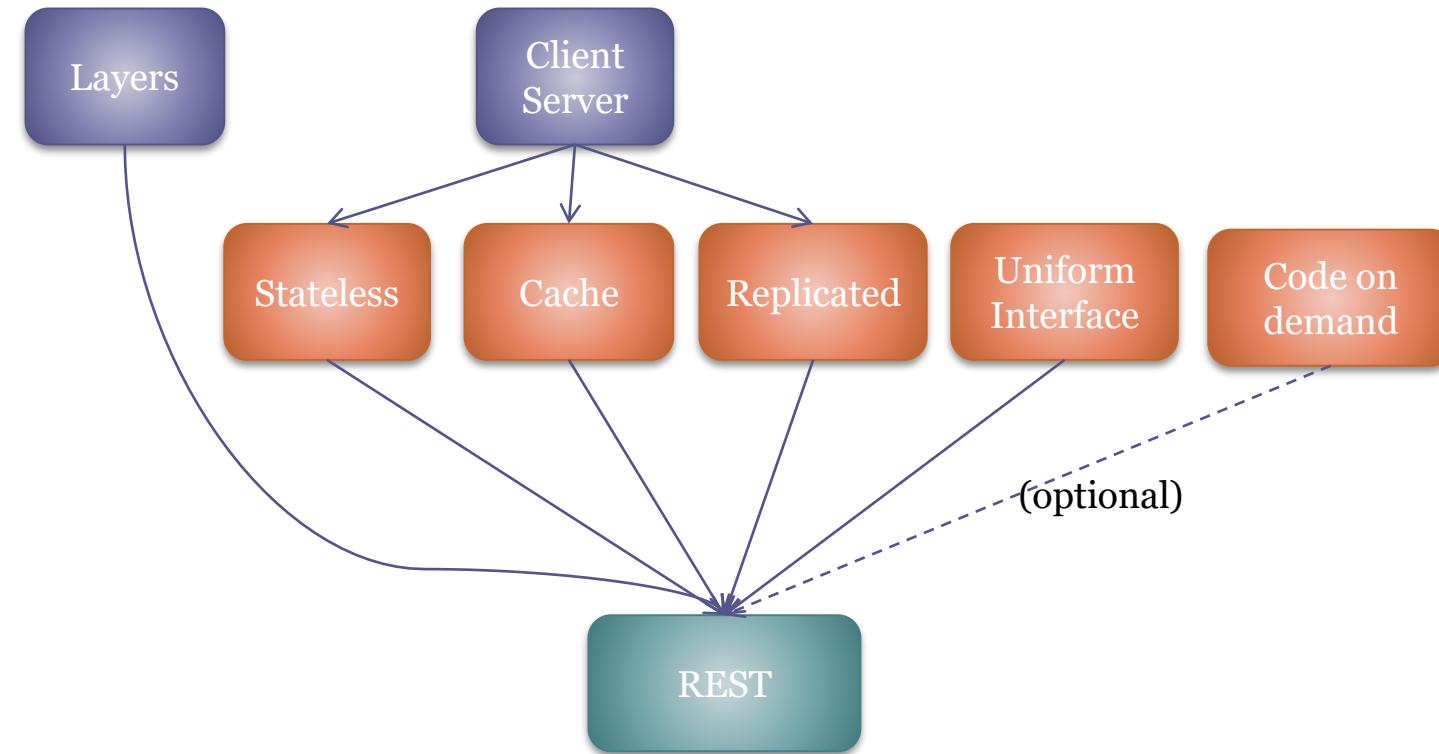
Web services without contract or description

RPC style REST

### Difficult to incorporate other requirements

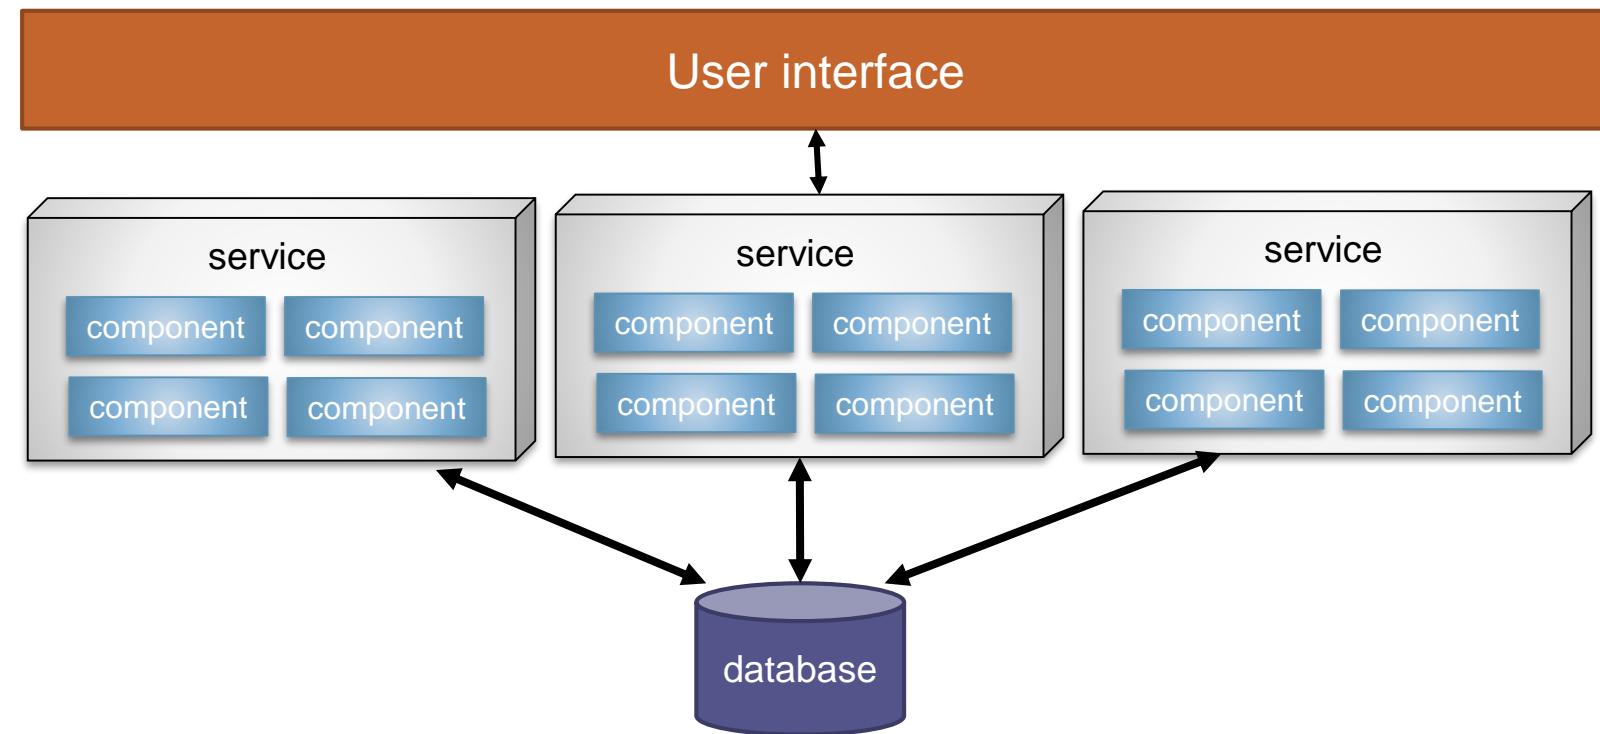
Security, transaction, composition, etc.

# REST as a composed style



# Service based architecture

Pragmatic architectural style based on SOA



# Service based architecture

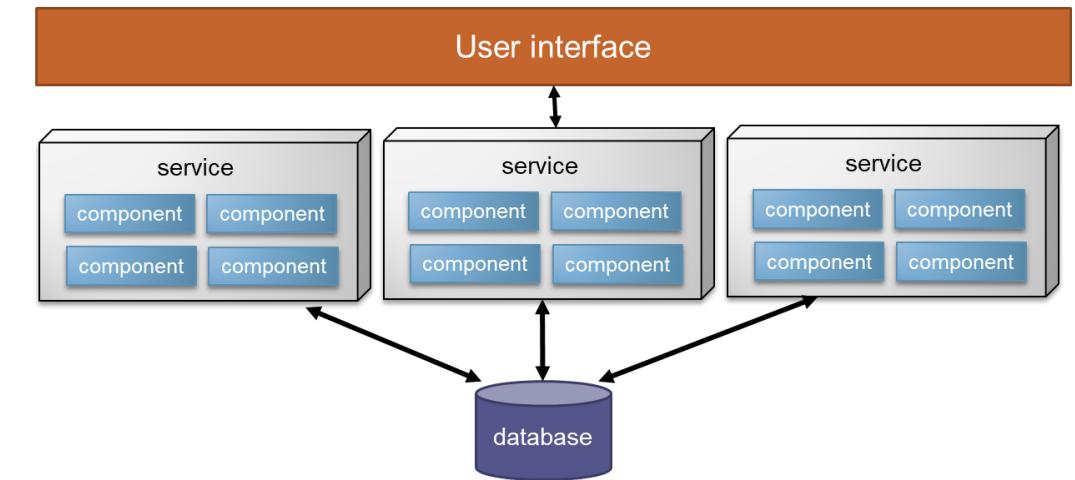
## Elements

**Services = independently deployed units**

Usually composed of different components

User interface accesses services remotely (Internet)

Database shared by those services



# Service based architecture

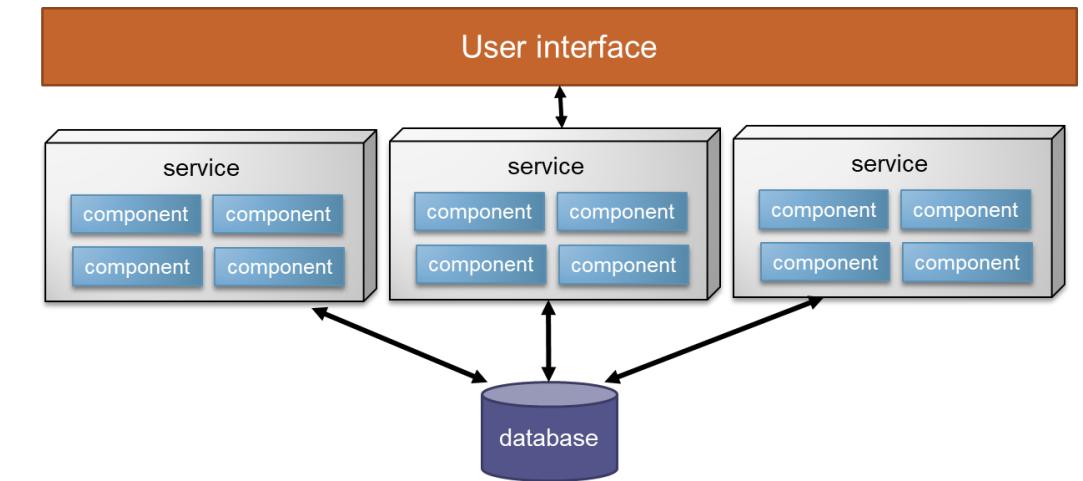
## Constraints

Each service is independently deployed

Services are usually coarse grained

User interface can be divided (different topologies)

Database is usually shared by each service



# Service based architecture

## Advantages

### Modularity of development

Services can be independently developed

### Technology diversity

Each service can be developed using a different programming language & technology

### Time to market

Several frameworks

### Availability

### Reliability

## Challenges

### Scalability (database partitioning)

### Evolution of services

Adaption to change is usually difficult

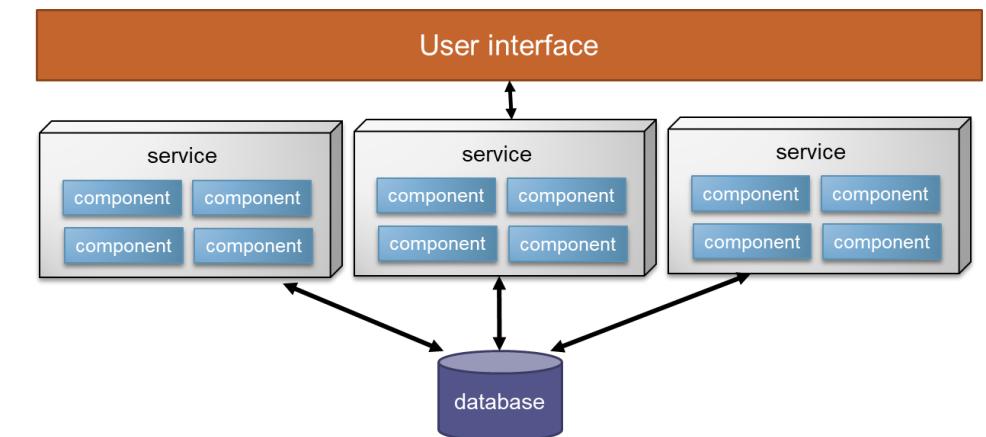
Services can be monoliths

### Conway's law

Database team

User interface team

Programmers



# Microservices

Applications decomposed in microservices

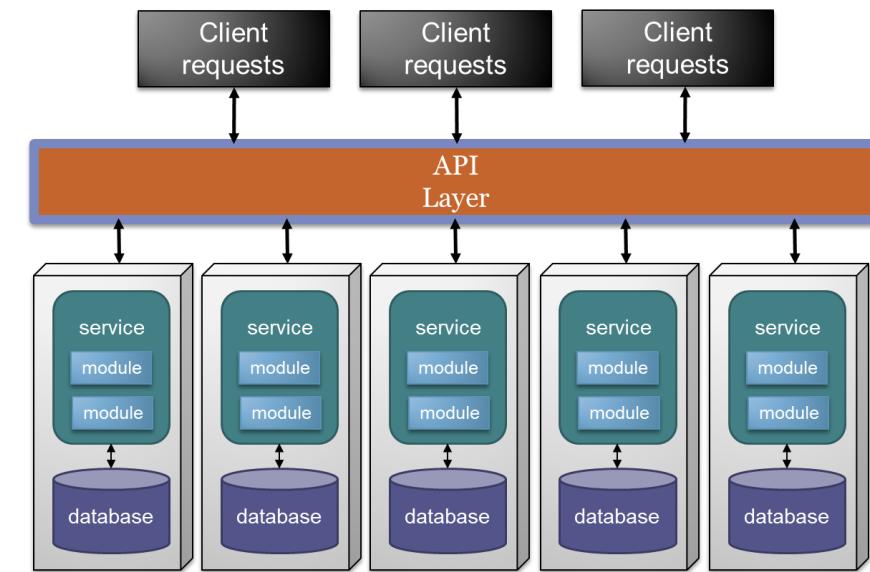
Microservice = small, autonomous services that work together

Each microservice = independent building and deployment block

Highly uncoupled

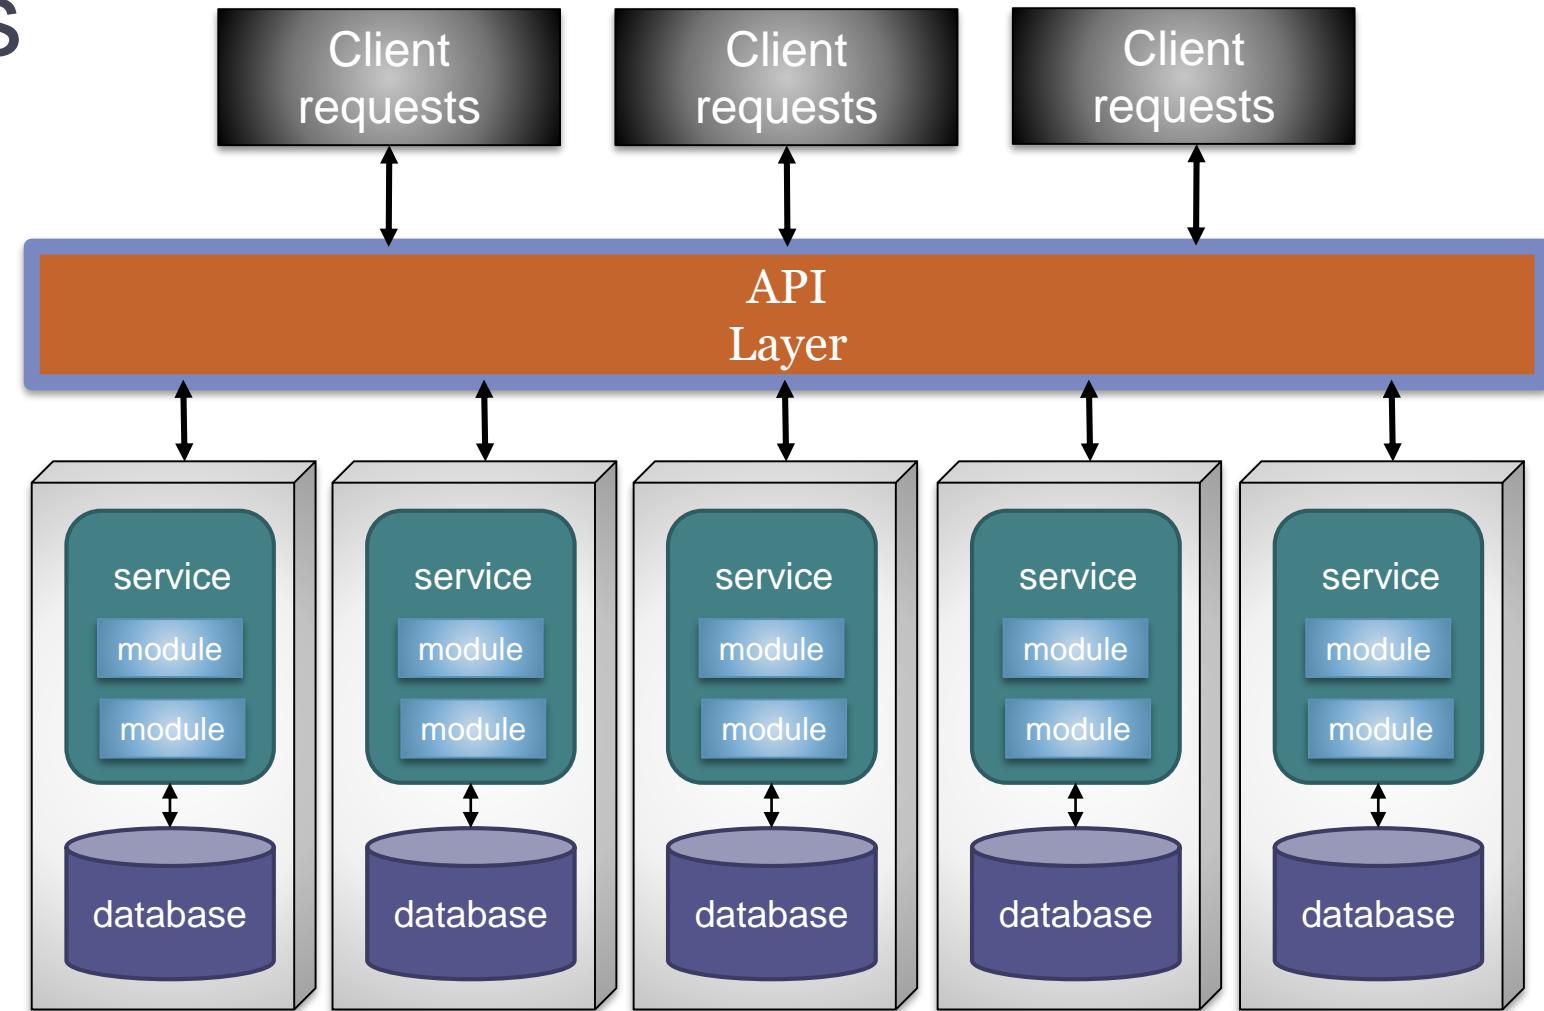
Focus on a specific task

Manage their own data



<http://martinfowler.com/articles/microservices.html>

# Microservices Diagram



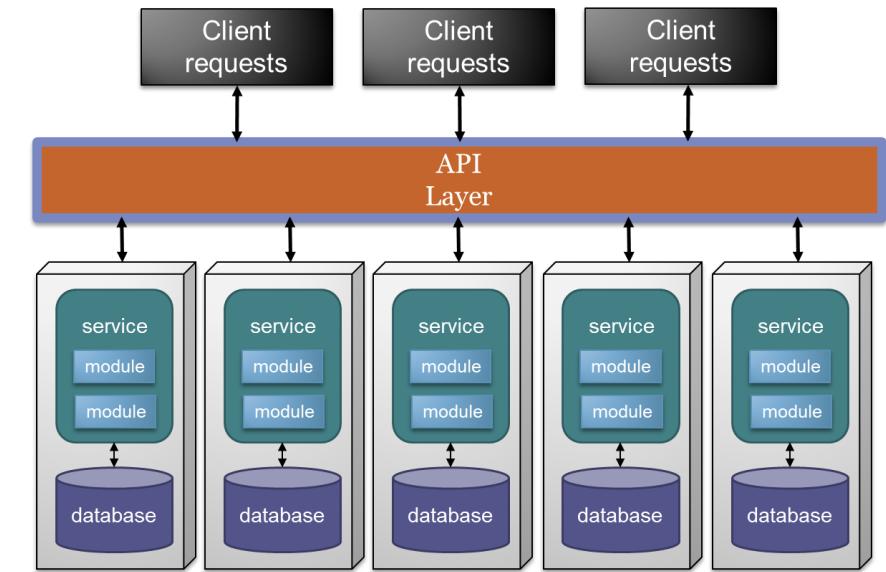
# Microservices

## Elements

A service + database form a deployed component

A service contains several modules and its own database

API layer (optional) offers a proxy or naming service



# Microservices

Constraints

Distributed

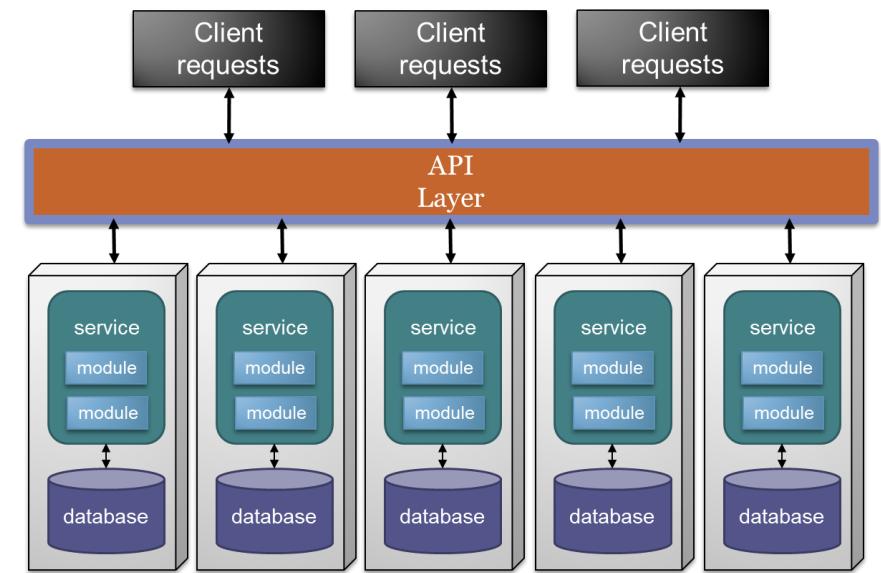
Bounded context:

Each service models a domain or workflow

Data isolation

Independency:

No mediator or orchestrator



# Features/advantages

Technology heterogeneity

Resilience

Scalability

Deployability

Organizational alignment

Decentralized data management

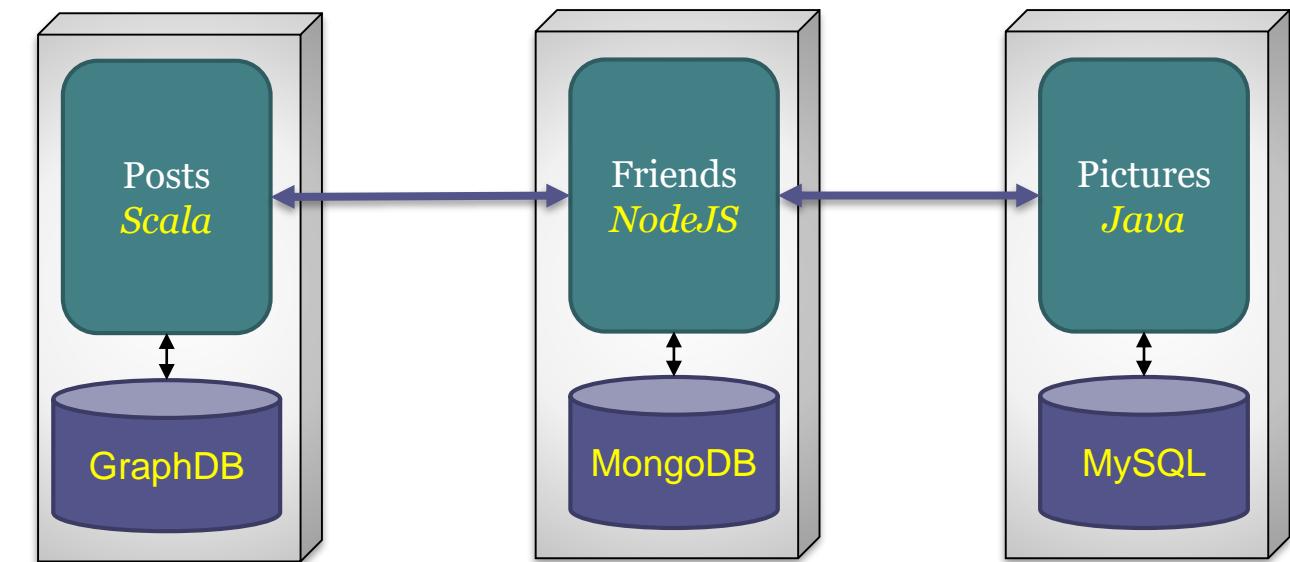
Optimizing for replaceability

# Microservices and heterogeneity

Each microservice can be implemented in its own programming language and technology stack

**Facilitates experimentation with new technologies**

**Flexibility**



# Resilience

If a component of a system fails and the failure doesn't scale, the system can carry on working

In a monolithic system if a component fails, the whole system stops working



# Scalability

It is possible to scale on demand specific services

Monolithic systems require to scale the whole system

Not all components have the same needs

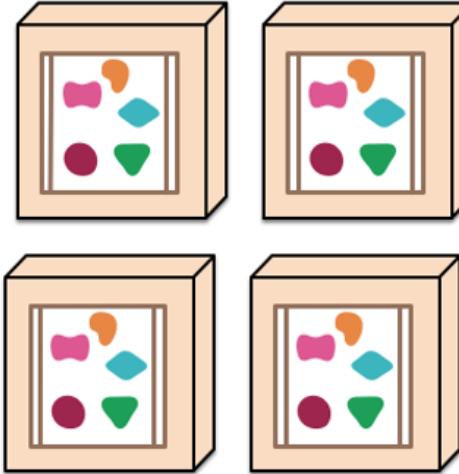
Microservices can be replicated as needed

# Scalability

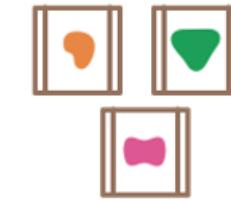
Monolithic: all functionality in a single process



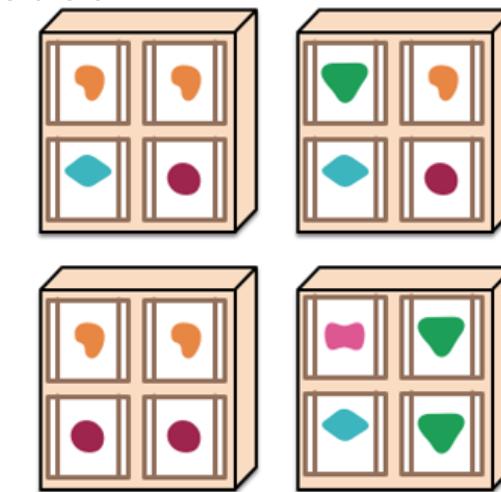
...scales replicating the monolith on multiple services



Microservices: each element of functionality into a separate service



... scales distributing these services and replicating as needed



# Deployability

Deploy each service independently

Enables to do a change in a service and deploy it immediately

Towards continuous deployment

# Organizational alignment

## Inverse Conway Law maneuver

Evolve teams and organizational structure to promote the desired architecture

Create teams following the modular decomposition

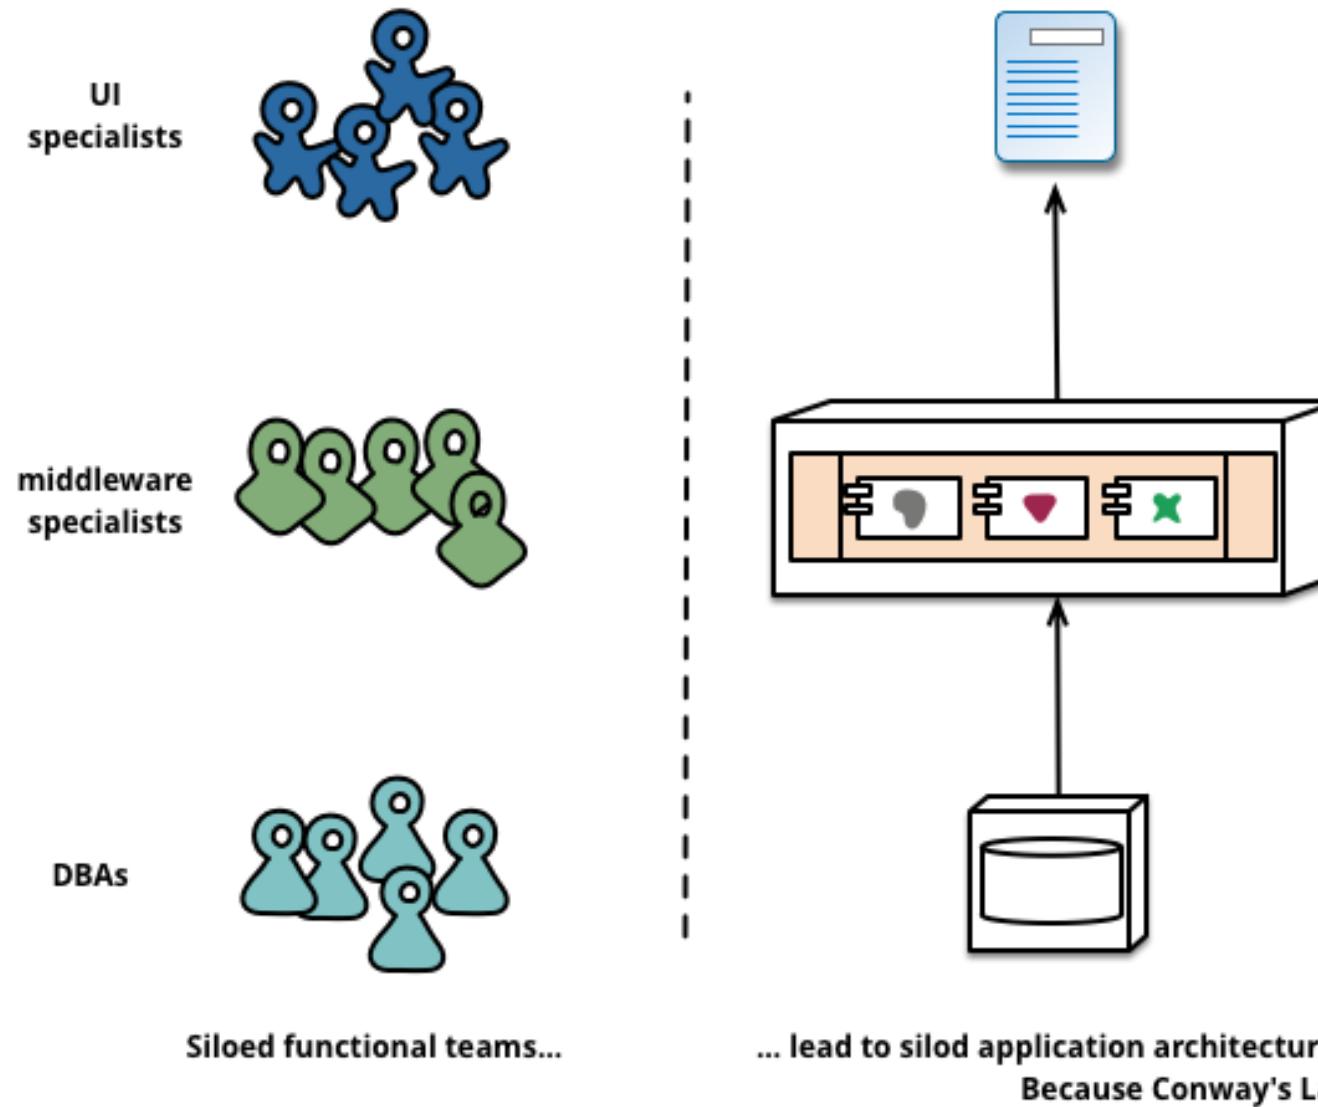
Cross-functional teams

Service ownership: the team owning a service is responsible for making changes and deploying it

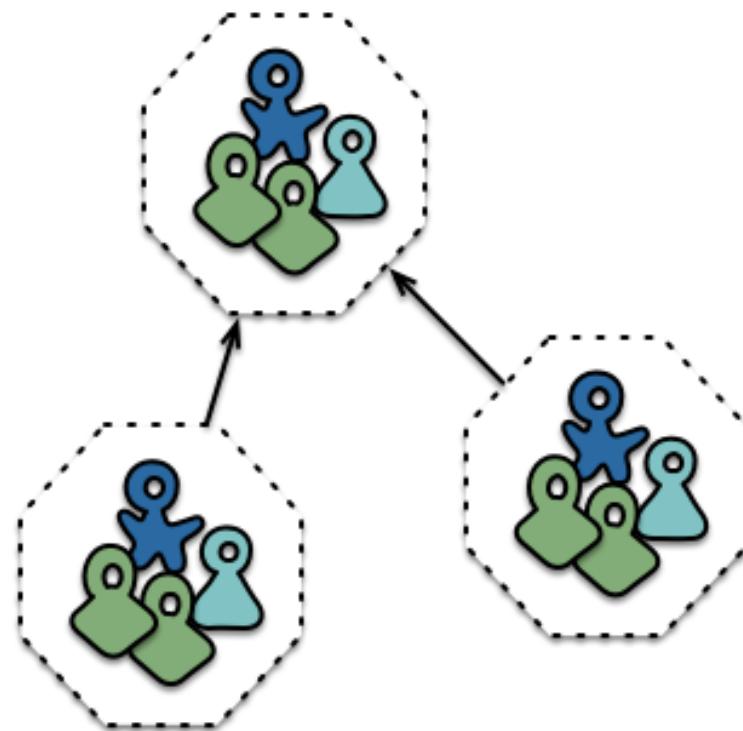
"You build it, you run it" (Amazon)

Goal: increased autonomy and speed of delivery

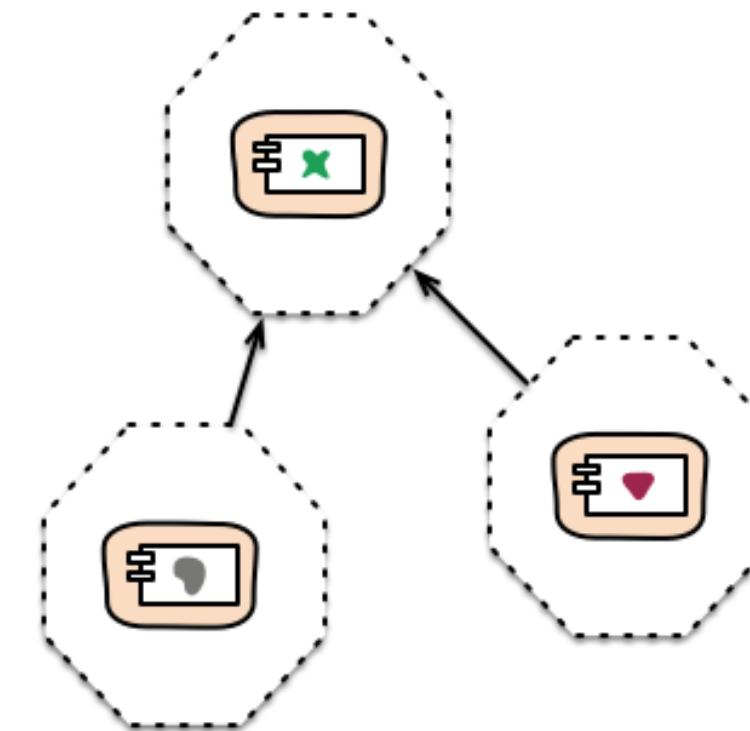
# Traditional applications



# With microservices



Cross-functional teams...

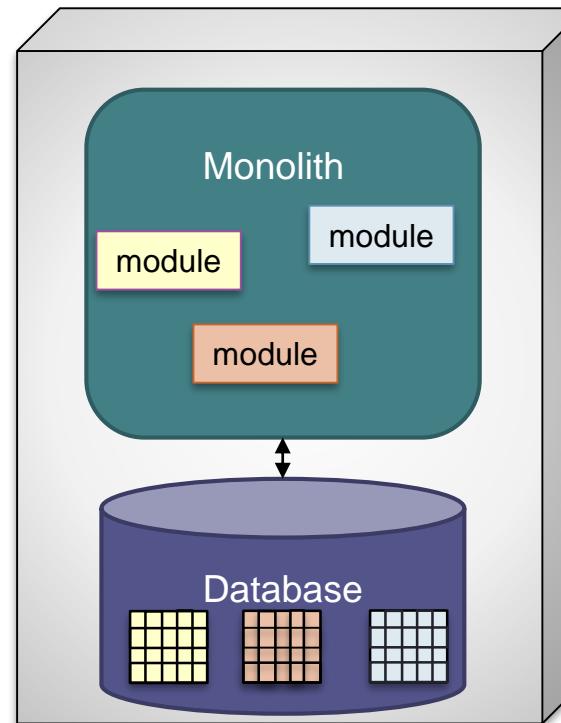


... organised around capabilities  
Because Conway's Law

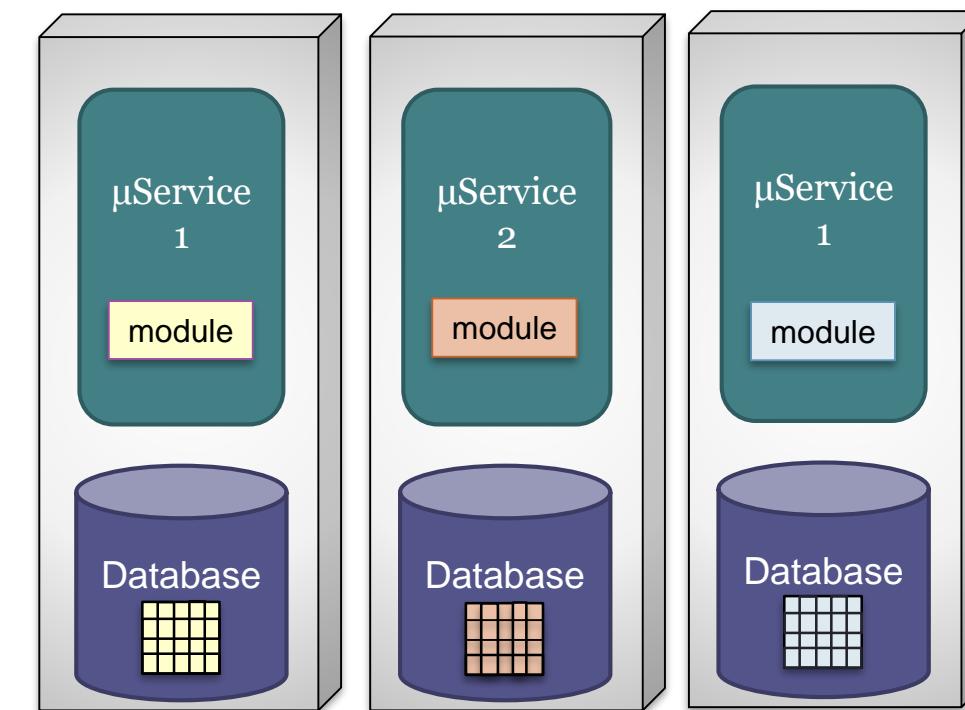
# Decentralized data management

Each team/service handles its own data

Monolith - single database



Microservices - application databases



# Optimizing for replaceability

Traditional systems usually contained old legacy systems which no one wants to touch

With microservices

Less cost to replace a microservice with a better implementation

Or even delete it

# Challenges of microservices

## Managing lots of microservices

Too much microservices = antipattern (nanoservices)

Ensure application consistency

## Complexity of distributed system management

New challenges: latency, message format, load balance, fault tolerance, etc.

## Testing & deployment

Operational complexity

## Antipattern: distributed monolith

Microservices tangled that are not independently deployed

## Structural decay (*see next slide*)

# Microservices structural decay

Code dependencies between services

Too much shared libraries

Too much inter-service communication

Too many orchestration requests

Database coupling

Analyzing architecture (microservices)  
<https://www.youtube.com/watch?v=U7s7Hb6GZCU>

# Microservices

## Variants

### **Self contained Systems (SCS) Architecture**

Separation of functionality into many independent systems

<https://scs-architecture.org/>

Each SCS contains logic and data

# Serverless

Also known as:

Function as a service (FaaS)

Backend as a service (BaaS)

Applications depend on third-party services

Developers don't need to care about servers

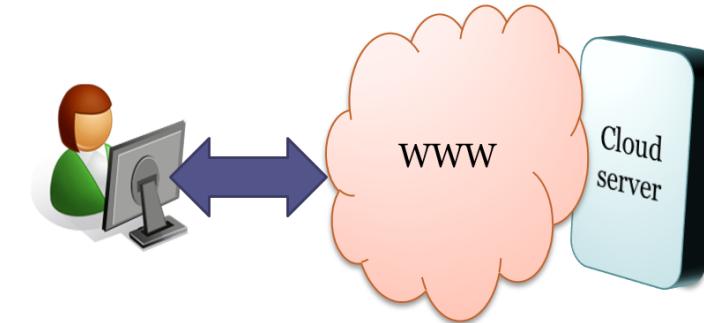
Automatic scalability

Rich clients

*Single Page Applications, Mobile apps*

Examples:

AWS Lambda, Google Cloud Functions, Ms Azure Functions



[https://en.wikipedia.org/wiki/Serverless\\_computing](https://en.wikipedia.org/wiki/Serverless_computing)  
<https://martinfowler.com/articles/serverless.html>

# Serverless

## Elements

Client that runs functions as a services

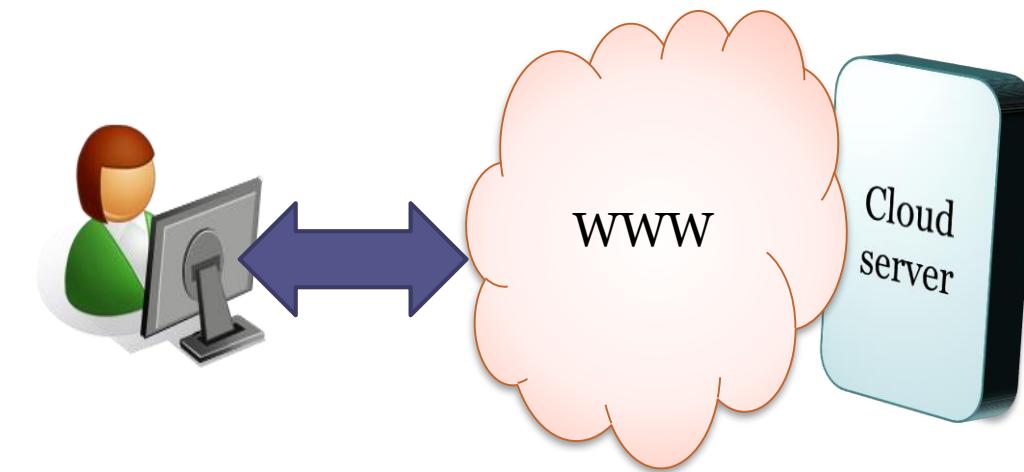
Cloud server which provides backend as a service

## Constraints

No management of server hosts

Automatic scalability and provisioning based on load

Costs based on precise usage



# Serverless

## Advantages

Automatic scalability

Implicit high availability

Performance not defined in terms of host size/cost

Costs based on precise usage

Only pay for the compute you need

Time to market

## Challenges

Vendor control

Vendor lock-in

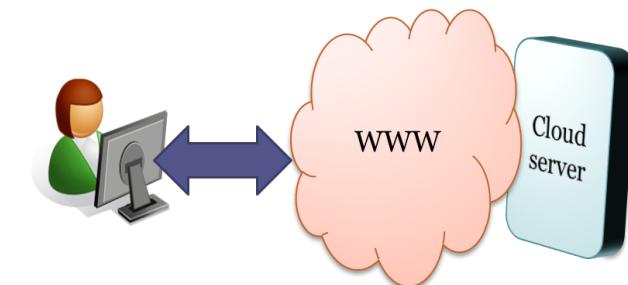
Incompatibility between vendors

Security

Startup latency

Integration testing

Monitoring/debugging



# Big data and scalable systems

MapReduce  
Lambda architecture  
Kappa architecture



# MapReduce

Proposed by Google

Published in 2004

Internal implementation by Google

Goal: big amounts of data

Lots of computational nodes

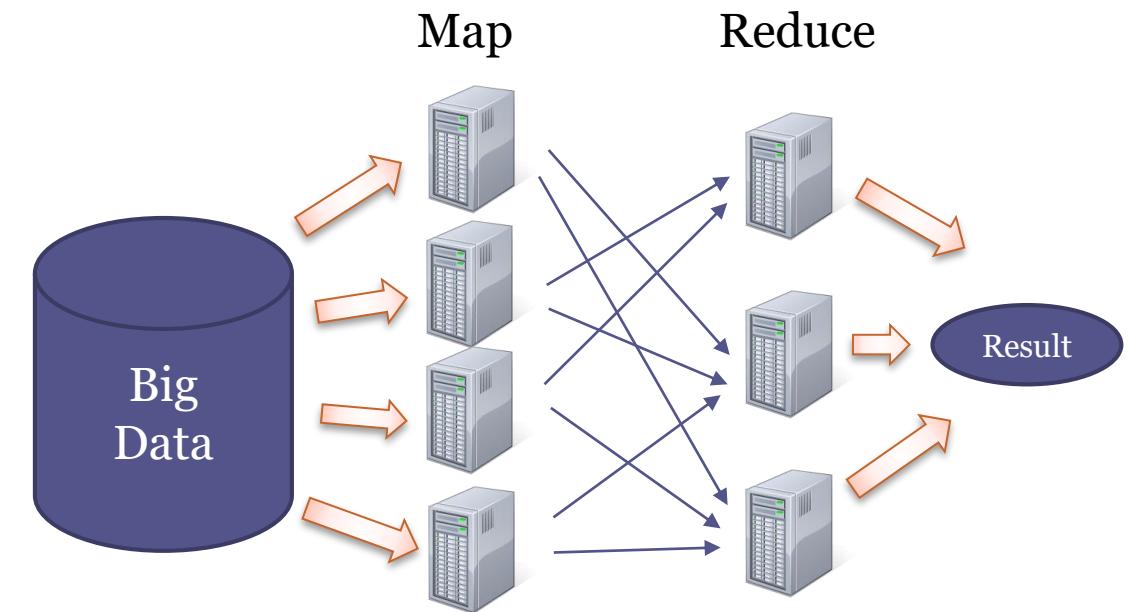
Fault tolerance

Write-once, read-many

Style composed of:

Master-slave

Batch



# MapReduce

## Elements

*Master node:* Controls execution

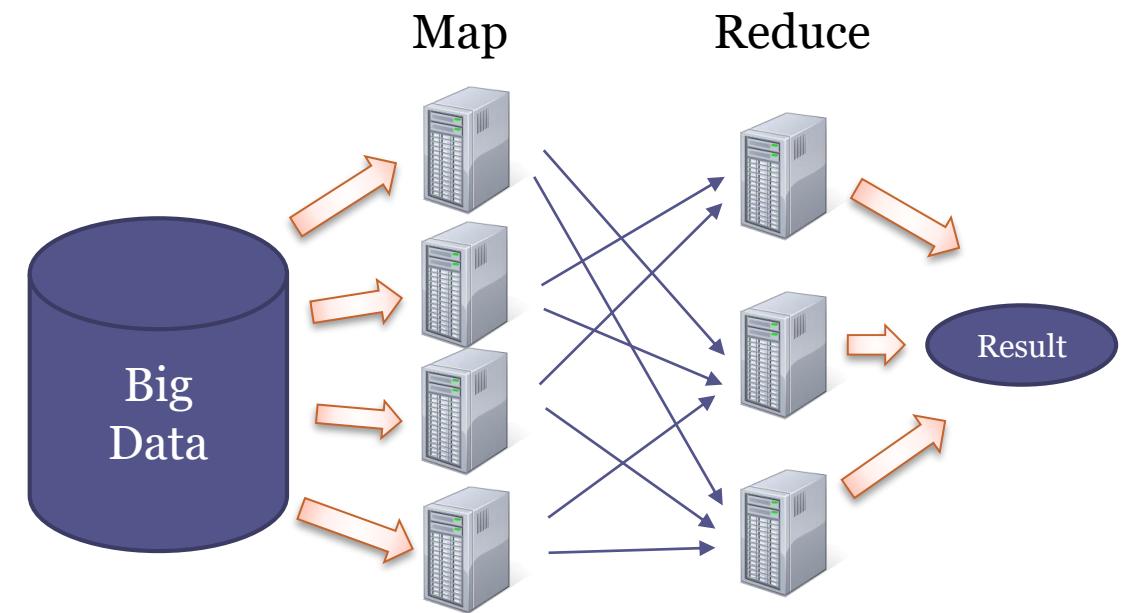
Node table

It manages replicated file system

*Slave nodes*

Execute mappers, reducers

Contain replicated data blocks



# MapReduce - Scheme

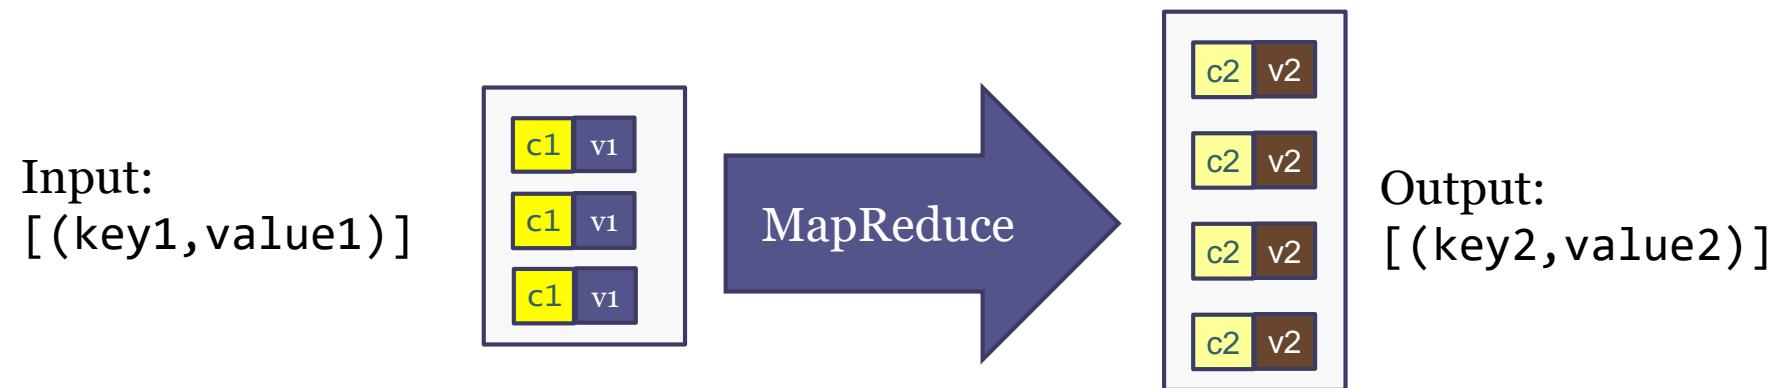
Inspired by functional programming

2 components: mapper and reducer

Data are divided for their processing

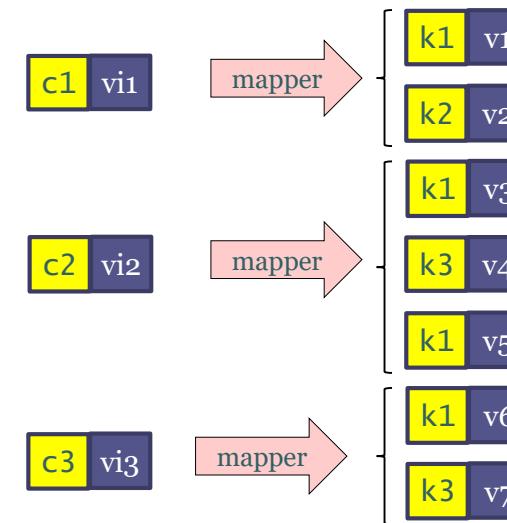
Each data is associated with a key

Transforms  $[(key1, value1)]$  to  $[(key2, value2)]$



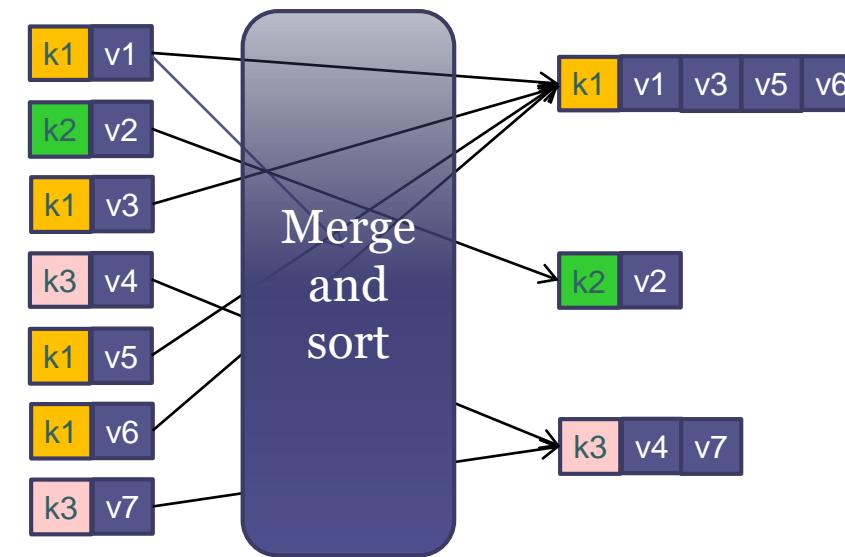
# Step 1: mapper

mapper: (Key1, Value1) → [(Key2,Value2)]



# Step 2: Merge and sort

System merges and sorts intermediate results according to the keys

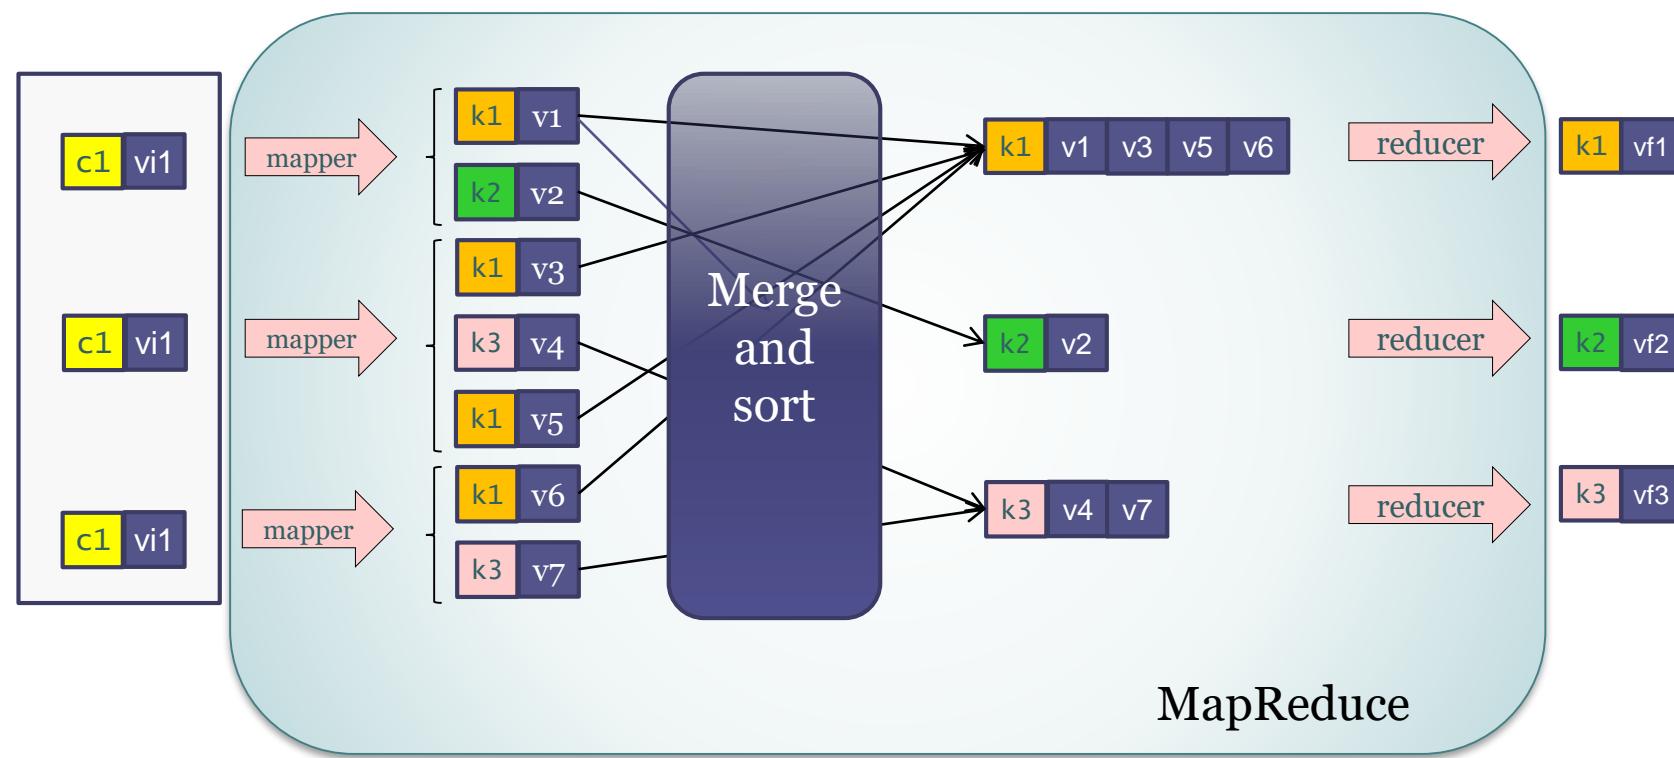


# Step 3: Reducers

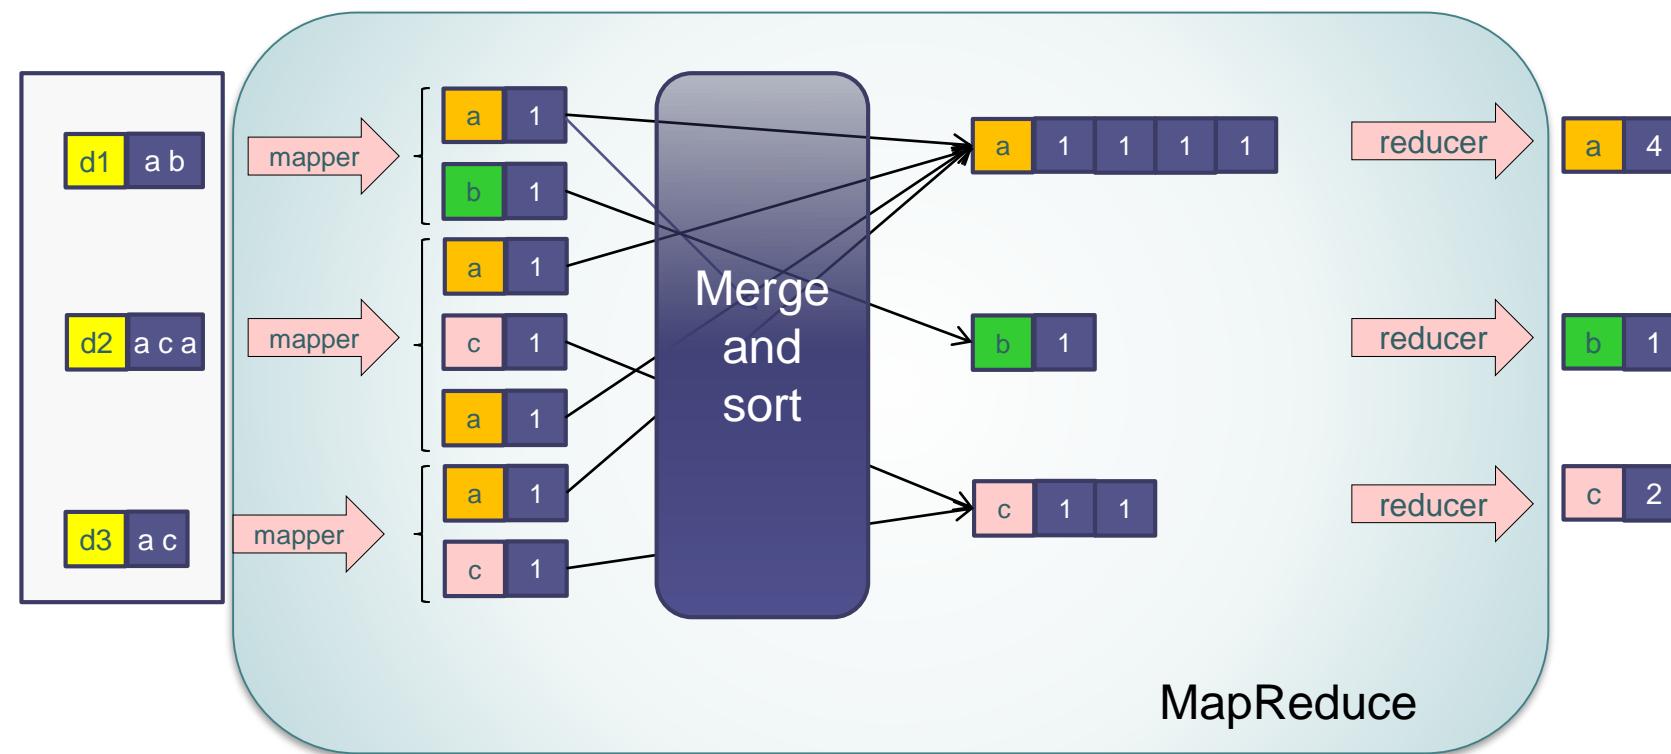
reducer: (Key2, [Value2]) → (Key2, Value2)



# MapReduce - general scheme



# MapReduce - count words



```
// return each word with 1
mapper(d,ps) {
    for each p in ps:
        emit (p, 1)
}
```

```
// sum the list of numbers of each word
reducer(p,ns) {
    sum = 0
    for each n in ns { sum += n; }
    emit (p, sum)
}
```

# MapReduce - execution environment

Execution environment is in charge of:

Planning: Each job is divided in tasks

Placement of data/code

Each node contains its data locally

Synchronization:

*reduce* tasks must wait *map* phase

Error and failure handling

High tolerance to computational nodes failures

# MapReduce - File system

Google developed a distributed file system - GFS

Hadoop created HDFS

Files are divided in chunks

2 node types:

Namenode (master), datanodes (data servers)

Datanodes store different chunks

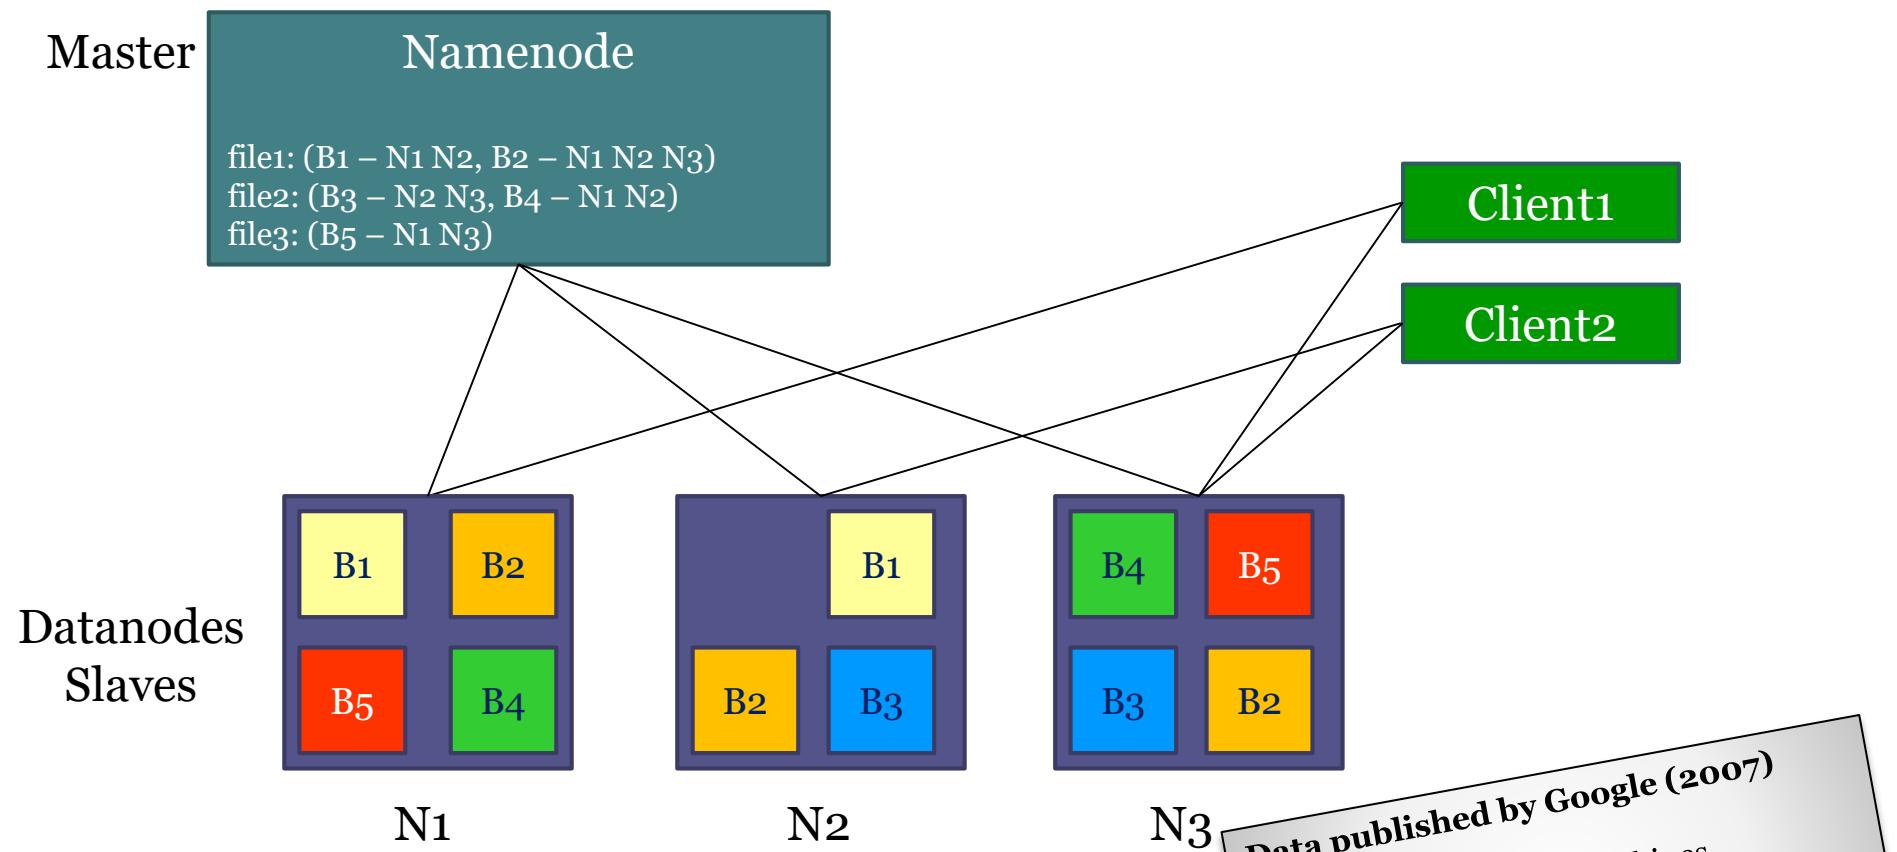
Block replication

*Namenode* contains metadata

Where is each chunk

Direct communication between clients and datanodes

# MapReduce - File system



**Data published by Google (2007)**

200+ clusters  
Lots of clusters 1000+ machines  
Pools with thousands of clients  
4+ PB  
HW/SW fault tolerance

# MapReduce

## Advantages

Distributed computations

Split input data

Replicated repository

Fault tolerant

Hardware/software  
heterogeneous

Large amount of data

Write-once. Read-many

## Challenges

Dependency on master node

Non interactivity

Data conversion to MapReduce

Adapt input data

Convert output data

# MapReduce: Applications

Lots of applications:

Google, 2007, 20petabytes/day, around 100,000 mapreduce jobs/day

PageRank algorithm can be implemented as MapReduce

Success stories:

Automatic translation, similarity, sorting, ...

Other companies: last.fm, facebook, Yahoo!, twitter, etc.

# MapReduce: Applications

## Implementations

Google (internal)

Hadoop (*open source*)

...

## Libraries

Hive (Hadoop): query language inspired by SQL

Pig (Hadoop): specific language that can define data flows

Cascading: API that can specify distributed data flows

Flume Java (Google)

Dryad (Microsoft)

# Lambda architecture



Handle Big Data & real time analytics

Proposed by Nathan Marz, 2011

3 layers

Batch layer: precomputes all data with MapReduce

Generates partial aggregate views

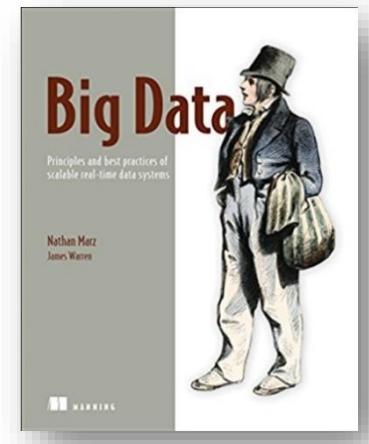
Recomputes from all data

Speed layer: real time, small window of data

Generates fast real time views

Serving layer: handles queries

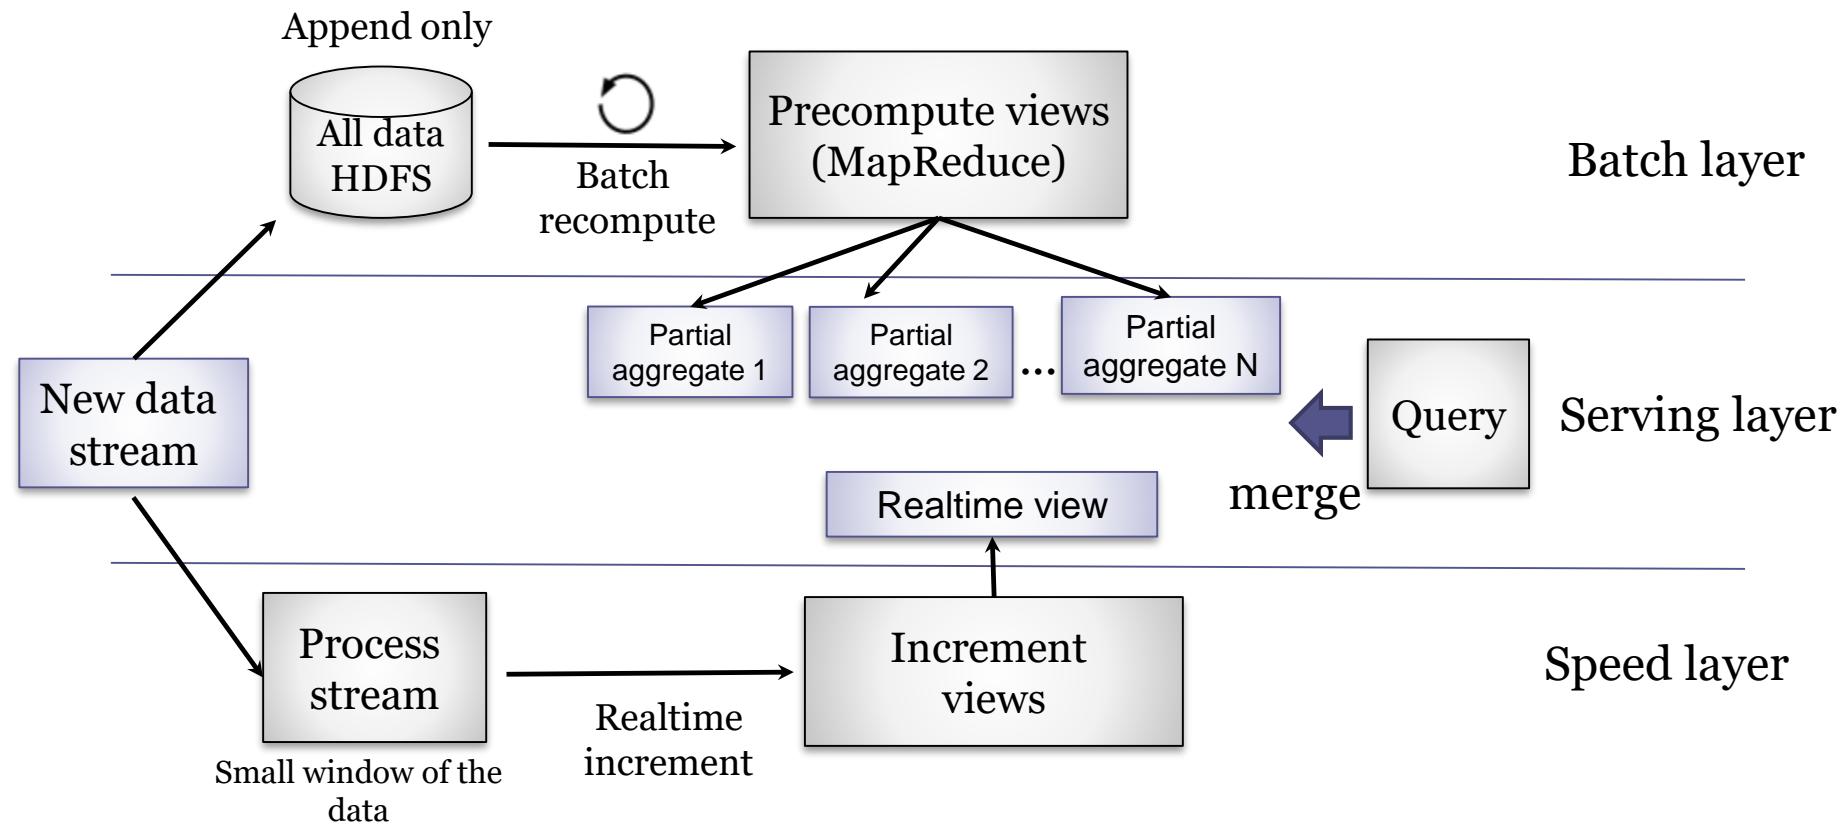
Merges the different views



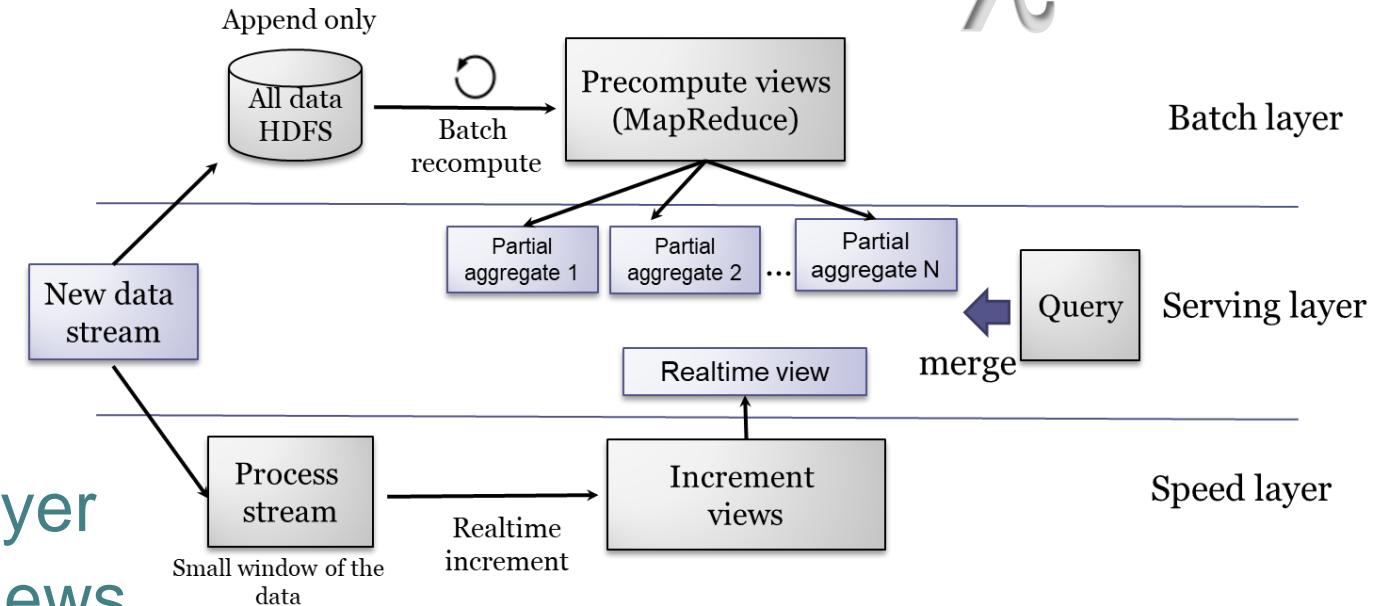
# Lambda architecture



Combines Real time with batch processing



# Lambda architecture



## Constraints

All data is stored in the batch layer

The batch layer precomputes views

The results of the speed layer may not be accurate

Serving layer combines precomputed views

The views can be simple DBs for querying

# Lambda architecture

 $\lambda$ 

## Advantages

Scalability (Big data)

Real time

Decoupling

Fault tolerant

Keep all input data

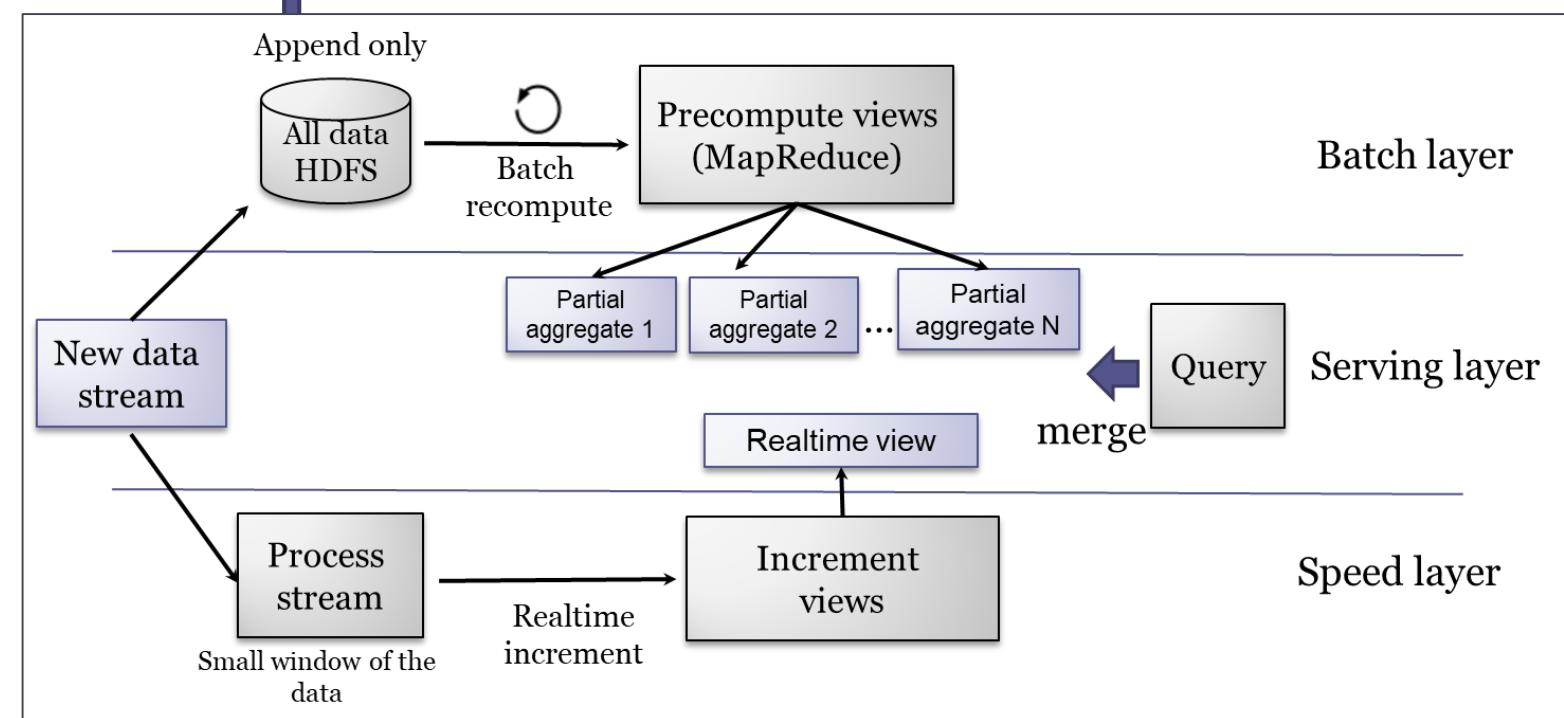
Reprocessing

## Challenges

Inherent complexity

Merging views can be inaccurate

Losing some events



# Lambda architecture



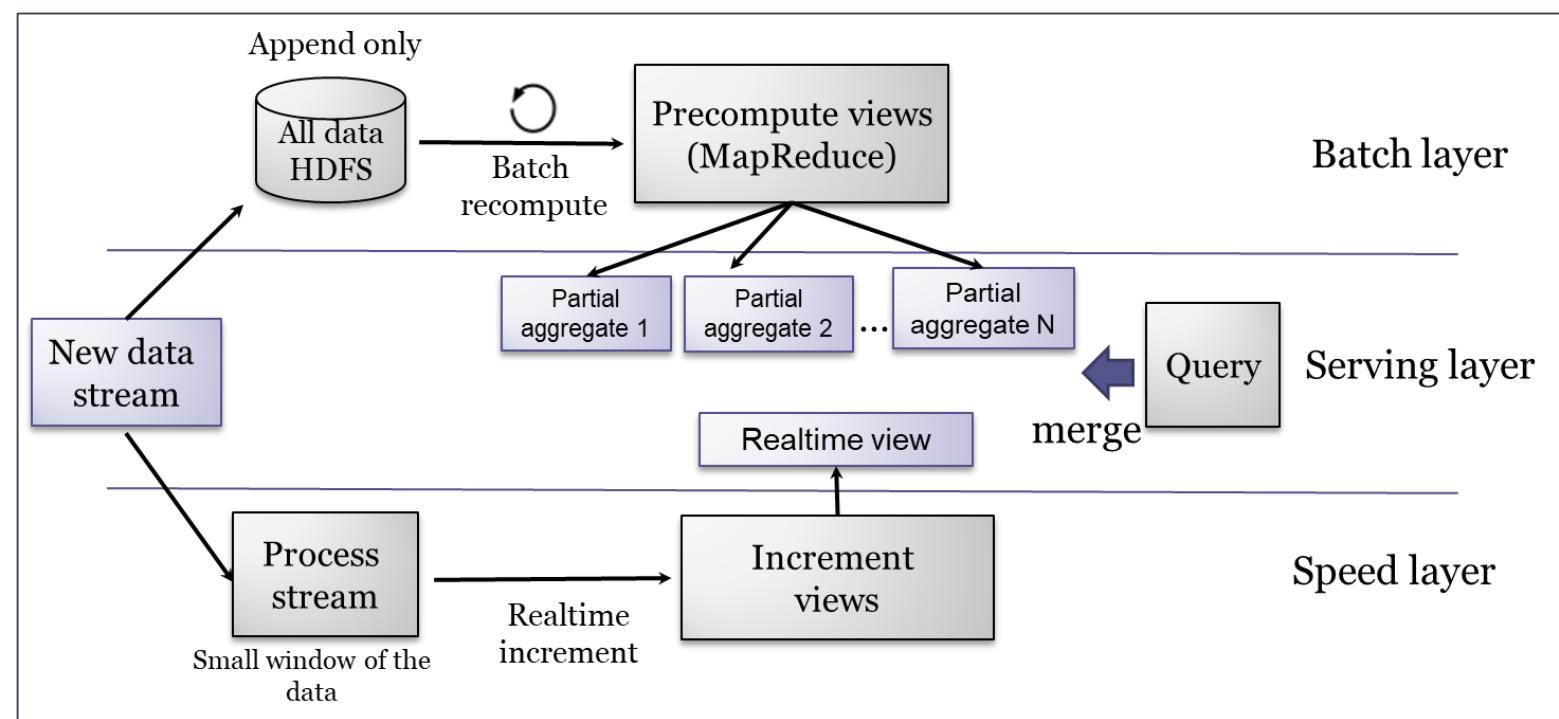
Applications

Spotify, Alibaba, ...

Libraries

Apache Storm

Netflix Suro project



# Kappa architecture



Proposed by Jay Krepps (Apache Kafka), 2013

Handle Big data & Real time with logs

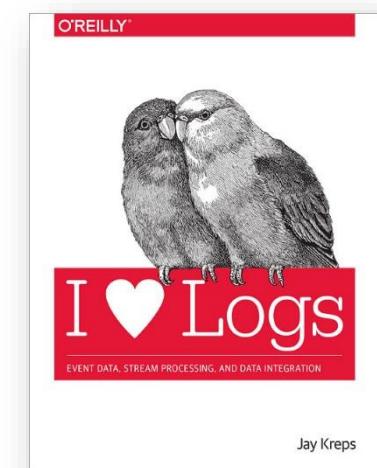
Simplifies Lambda architecture

Removes the batch layer

Based on a distributed ordered log

Replicated cluster

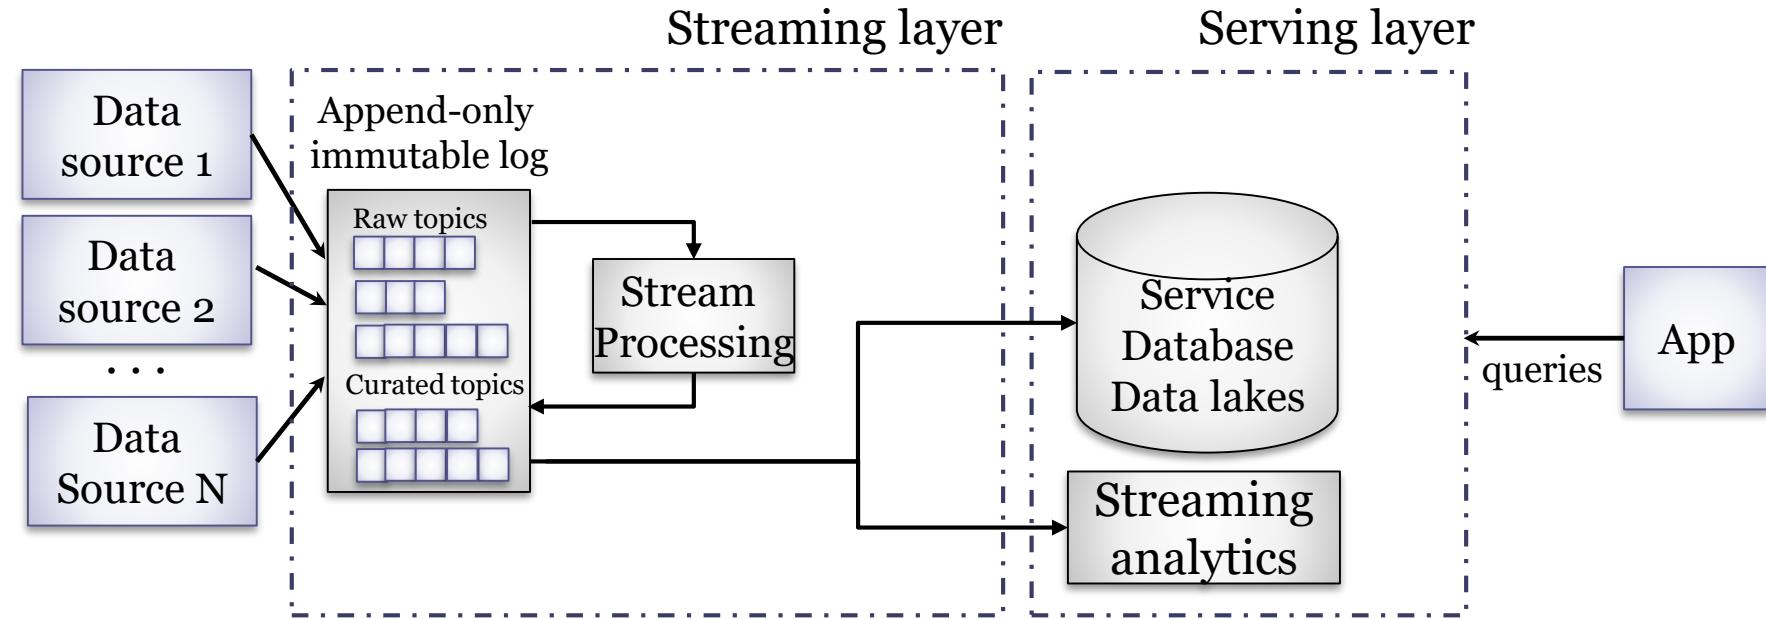
The log can be very large



# Kappa architecture



## Diagram



# Kappa architecture



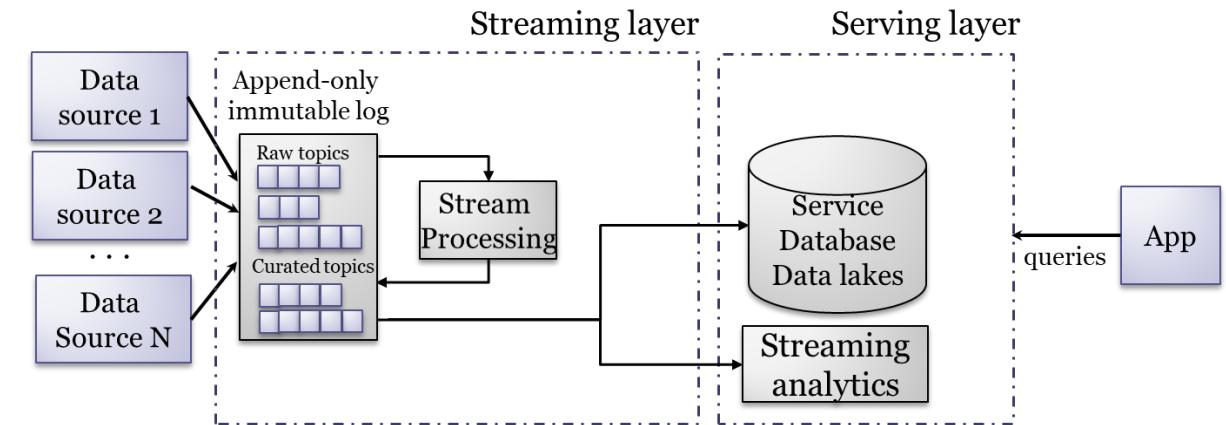
## Constraints

The event log is append-only

The events in the log are immutable

Stream processing can request events at any position

To handle failures or doing recomputations



# Kappa architecture



## Advantages

- Scalable (big data)
- Real time processing
- Simpler than lambda
- No batch layer

## Challenges

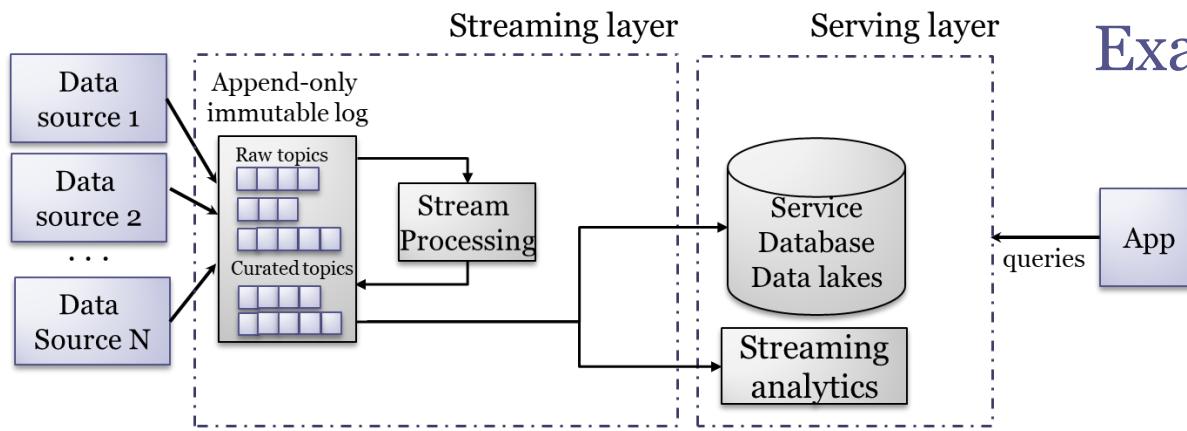
### Space requirements

- Duplication of log and DB
- Log compaction

### Ordering of events

### Delivery paradigms

- At least once, at most once (it may be lost),  
Exactly once



# Kappa architecture



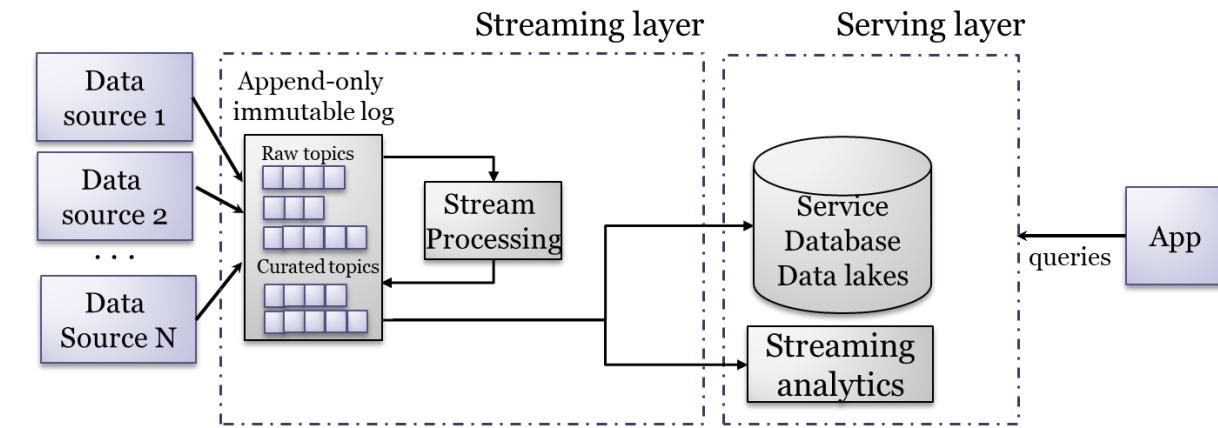
## Applications & libraries

Apache Kafka

Apache Samza

Spark Streaming

LinkedIn



**End of presentation**